

# ADS 503 - Team 7

Summer Purschke, Jacqueline Urenda, Oscar Gil

06/12/2022

```
# R Libraries
library(caret)
library(AppliedPredictiveModeling)
library(Hmisc)
library(dplyr)
library(tidyverse)
library(ggplot2)
library(corrplot)
library(MASS)
library(ISLR)
library(rpart)
library(partykit)
library(randomForestSRC)
library(earth)
library(MARSS)
library(e1071)
library(summarytools)
```

**Load the Red Wine Quality data set from GitHub - data set copied from Kaggle and imported into GitHub.**

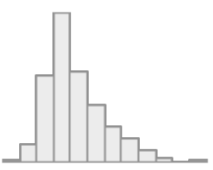
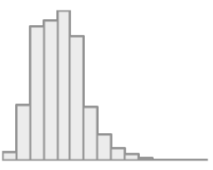
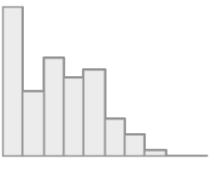
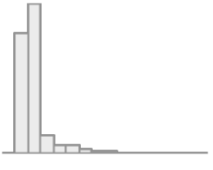
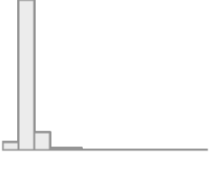
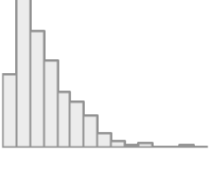
```
wine <- read.csv(
  url("https://raw.githubusercontent.com/OscarG-DataSci/ADS503/main/winequality-red.csv"),
  , header = TRUE)
```

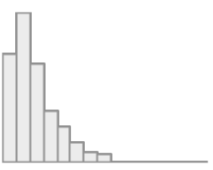
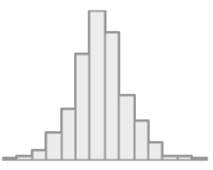
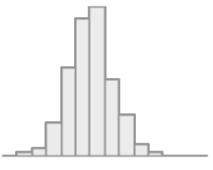
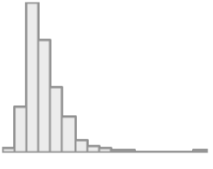
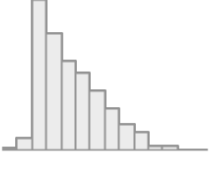
## Data Summary


```
dfSummary(wine,
  plain.ascii = FALSE,
  style       = "grid",
  graph.magnif = 0.75,
  valid.col    = FALSE,
  tmp.img.dir  = "/tmp")
```

## Data Frame Summary

wine Dimensions: 1599 x 12  
Duplicates: 240

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
1	fixed.acidity [numeric]	Mean (sd) : 8.3 (1.7) min < med < max: 4.6 < 7.9 < 15.9 IQR (CV) : 2.1 (0.2)	96 distinct values		0 (0.0%)
2	volatile.acidity [numeric]	Mean (sd) : 0.5 (0.2) min < med < max: 0.1 < 0.5 < 1.6 IQR (CV) : 0.2 (0.3)	143 distinct values		0 (0.0%)
3	citric.acid [numeric]	Mean (sd) : 0.3 (0.2) min < med < max: 0 < 0.3 < 1 IQR (CV) : 0.3 (0.7)	80 distinct values		0 (0.0%)
4	residual.sugar [numeric]	Mean (sd) : 2.5 (1.4) min < med < max: 0.9 < 2.2 < 15.5 IQR (CV) : 0.7 (0.6)	91 distinct values		0 (0.0%)
5	chlorides [numeric]	Mean (sd) : 0.1 (0) min < med < max: 0 < 0.1 < 0.6 IQR (CV) : 0 (0.5)	153 distinct values		0 (0.0%)
6	free.sulfur.dioxide [numeric]	Mean (sd) : 15.9 (10.5) min < med < max: 1 < 14 < 72 IQR (CV) : 14 (0.7)	60 distinct values		0 (0.0%)

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
7	total.sulfur.dioxide [numeric]	Mean (sd) : 46.5 (32.9) min < med < max: 6 < 38 < 289 IQR (CV) : 40 (0.7)	144 distinct values		0 (0.0%)
8	density [numeric]	Mean (sd) : 1 (0) min < med < max: 1 < 1 < 1 IQR (CV) : 0 (0)	436 distinct values		0 (0.0%)
9	pH [numeric]	Mean (sd) : 3.3 (0.2) min < med < max: 2.7 < 3.3 < 4 IQR (CV) : 0.2 (0)	89 distinct values		0 (0.0%)
10	sulphates [numeric]	Mean (sd) : 0.7 (0.2) min < med < max: 0.3 < 0.6 < 2 IQR (CV) : 0.2 (0.3)	96 distinct values		0 (0.0%)
11	alcohol [numeric]	Mean (sd) : 10.4 (1.1) min < med < max: 8.4 < 10.2 < 14.9 IQR (CV) : 1.6 (0.1)	65 distinct values		0 (0.0%)

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
					
12	quality [integer]	Mean (sd) : 5.6 (0.8) min < med < max: 3 < 6 < 8 IQR (CV) : 1 (0.1)	3 : 10 ( 0.6%) 4 : 53 ( 3.3%) 5 : 681 (42.6%) 6 : 638 (39.9%) 7 : 199 (12.4%) 8 : 18 ( 1.1%)		0 (0.0%)

## Pre-processing

```
par(mar=c(1,1,1,1)) # to fix boxplot knit processing issues
```

```
# Create new variable, for quality values, split by half (0, 1)
wine$quality_target <- ifelse( wine$quality <= 5, 0, 1)
```

```
# Mean of new variable is at 0.5347 (close enough to 50% to maintain balance)
summary(wine$quality_target)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000  0.0000  1.0000  0.5347  1.0000  1.0000
```

```
# Check for missing values in data set
wine %>% na.omit() %>% count() # there are no missing values
```

```
##      n
## 1 1599
```

```
# Removing outliers for residual sugar:
```

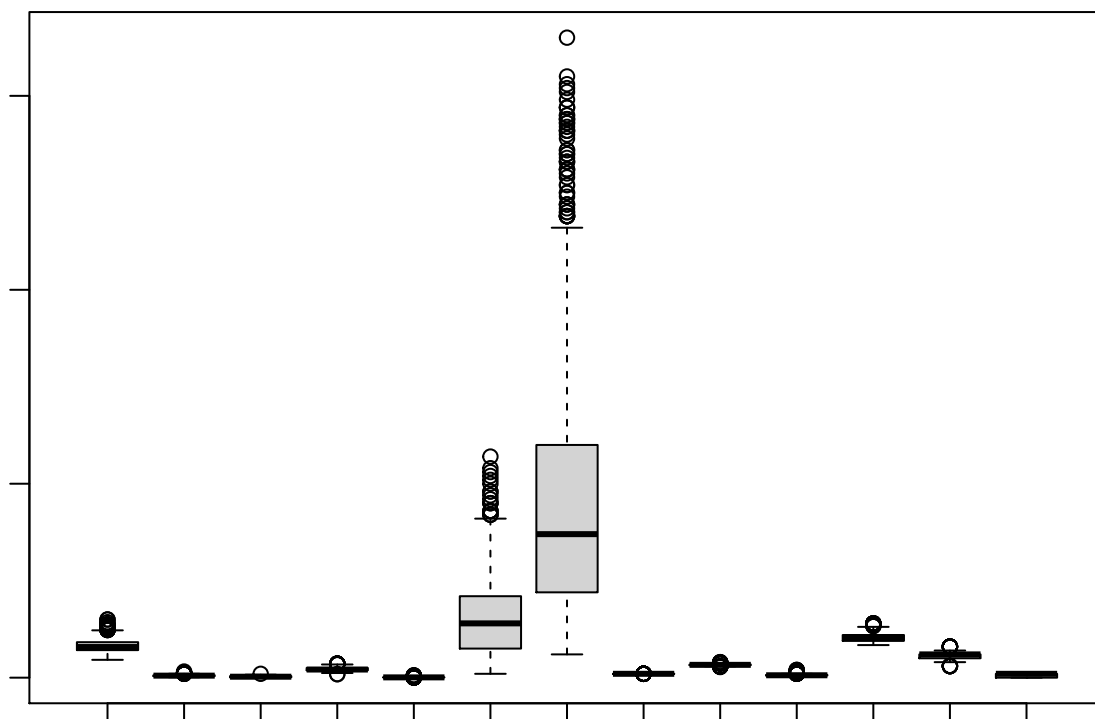
```
Q <- quantile(wine$residual.sugar, probs=c(.25, .75), na.rm = FALSE)
```

```
iqr_rs <- IQR(wine$residual.sugar)
```

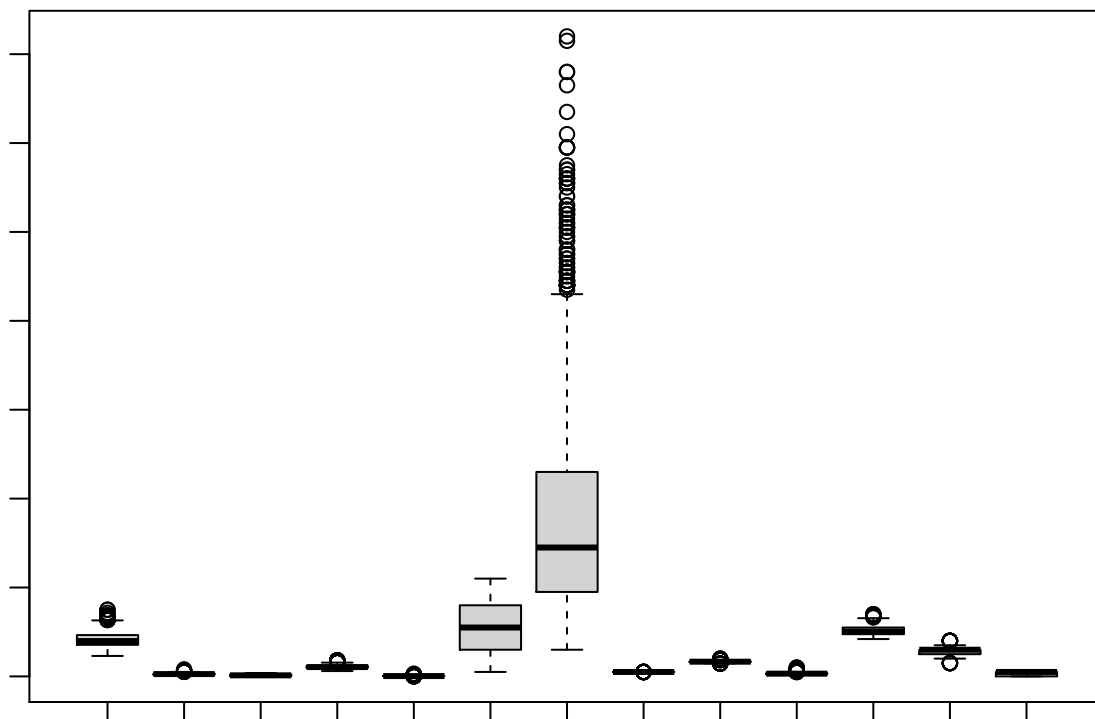
```
up_rs <- Q[2]+1.5*iqr_rs # Upper Range
```

```
low_rs <- Q[1]-1.5*iqr_rs # Lower Range
```

```
eliminated_rs <- subset(wine, wine$residual.sugar > (Q[1] - 1.5*iqr_rs) & wine$residual.sugar < (Q[2]+1.5*iqr_rs))
boxplot(eliminated_rs)
```



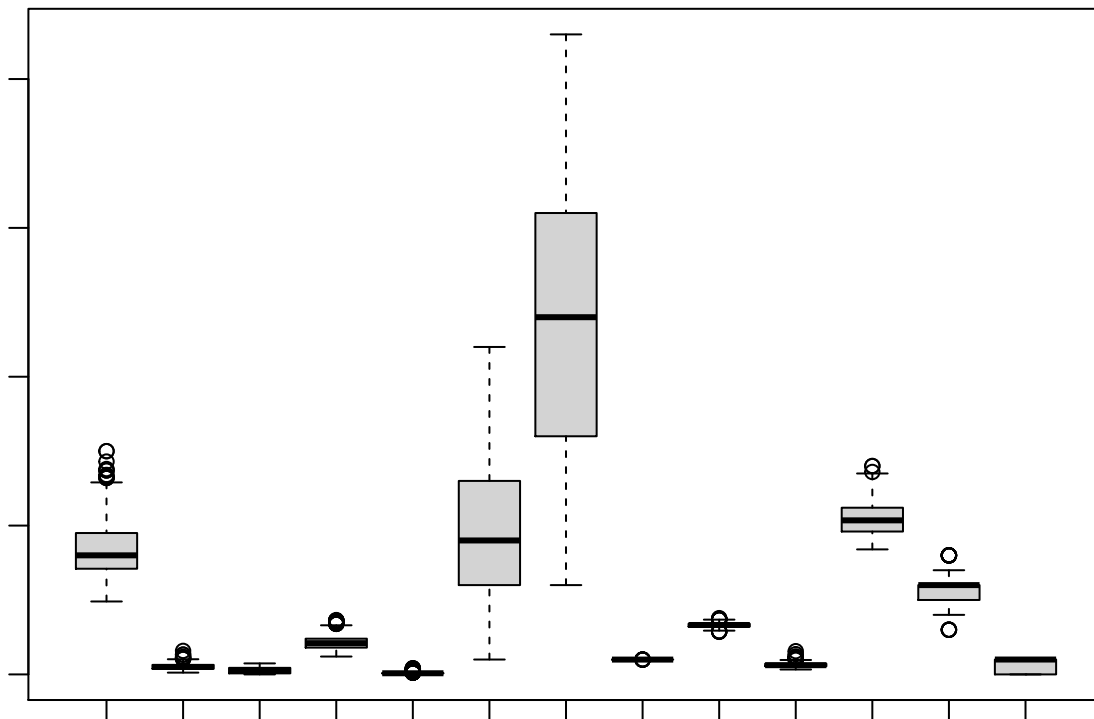
```
#Removing outliers for free.sulfur.dioxide:
Q2 <- quantile(wine$free.sulfur.dioxide, probs=c(.25, .75), na.rm = FALSE)
iqr_fs <- IQR(eliminated_rs$free.sulfur.dioxide)
up_fs <- Q2[2]+1.5*iqr_fs # Upper Range
low_fs <- Q2[1]-1.5*iqr_fs # Lower Range
eliminated_fs <- subset(eliminated_rs, eliminated_rs$free.sulfur.dioxide > (Q[1] - 1.5*iqr_fs) & eliminated_rs$free.sulfur.dioxide < (Q[2] + 1.5*iqr_fs))
boxplot(eliminated_fs)
```



```

#Removing outliers for total.sulfur.dioxide:
Q3 <- quantile(wine$total.sulfur.dioxide, probs=c(.25, .75), na.rm = FALSE)
iqr_ts <- IQR(eliminated_fs$total.sulfur.dioxide)
up_ts <- Q3[2]+1.5*iqr_ts # Upper Range
low_ts <- Q3[1]-1.5*iqr_ts # Lower Range
eliminated_ts <- subset(eliminated_fs, eliminated_fs$total.sulfur.dioxide > (Q[1] - 1.5*iqr_ts) & eliminated_ts < up_ts)
boxplot(eliminated_ts)

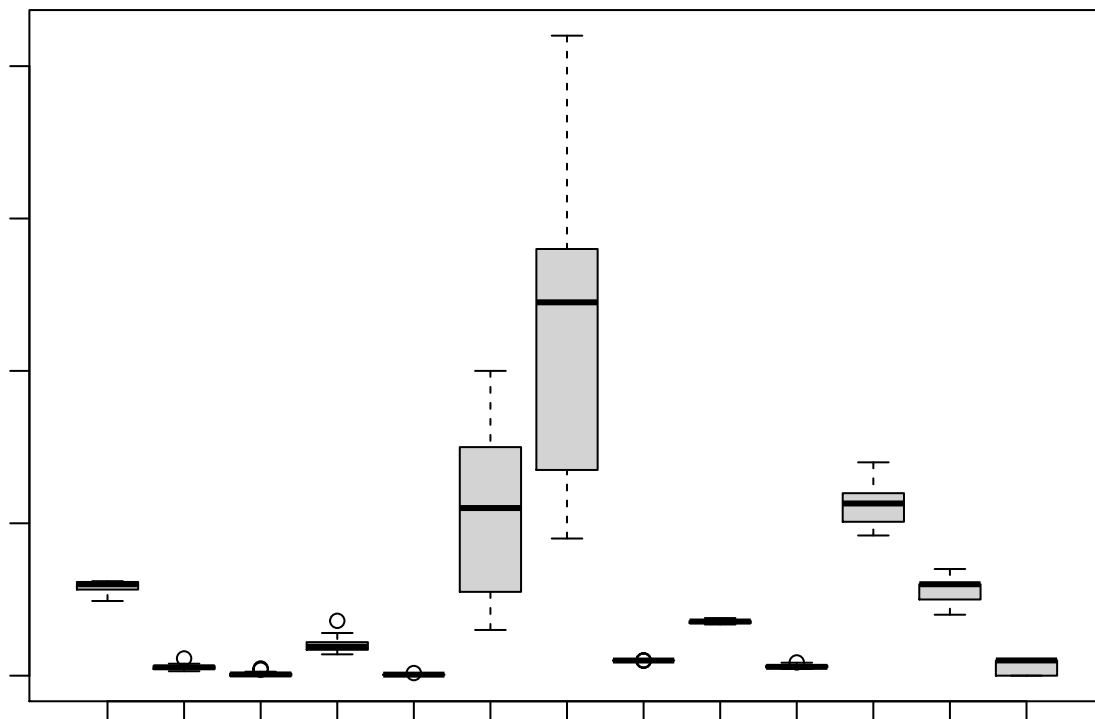
```



```

#Removing outliers for fixed.acidity:
Q4 <- quantile(wine$fixed.acidity, probs=c(.25, .75), na.rm = FALSE)
iqr_fa <- IQR(eliminated_ts$fixed.acidity)
up_fa <- Q[2]+1.5*iqr_fa # Upper Range
low_fa <- Q[1]-1.5*iqr_fa # Lower Range
eliminated_fa <- subset(eliminated_ts, eliminated_ts$fixed.acidity > (Q[1] - 1.5*iqr_fa) & eliminated_ts < up_fa)
boxplot(eliminated_fa)

```



```
new_wine_data <- eliminated_fa
```

```
# Removing outliers reduced dimension of data set from 1599 observations to 48
```

```
# team opted not to use new_wine_data and keep outlier data
```

```
dim(new_wine_data)
```

```
## [1] 48 13
```

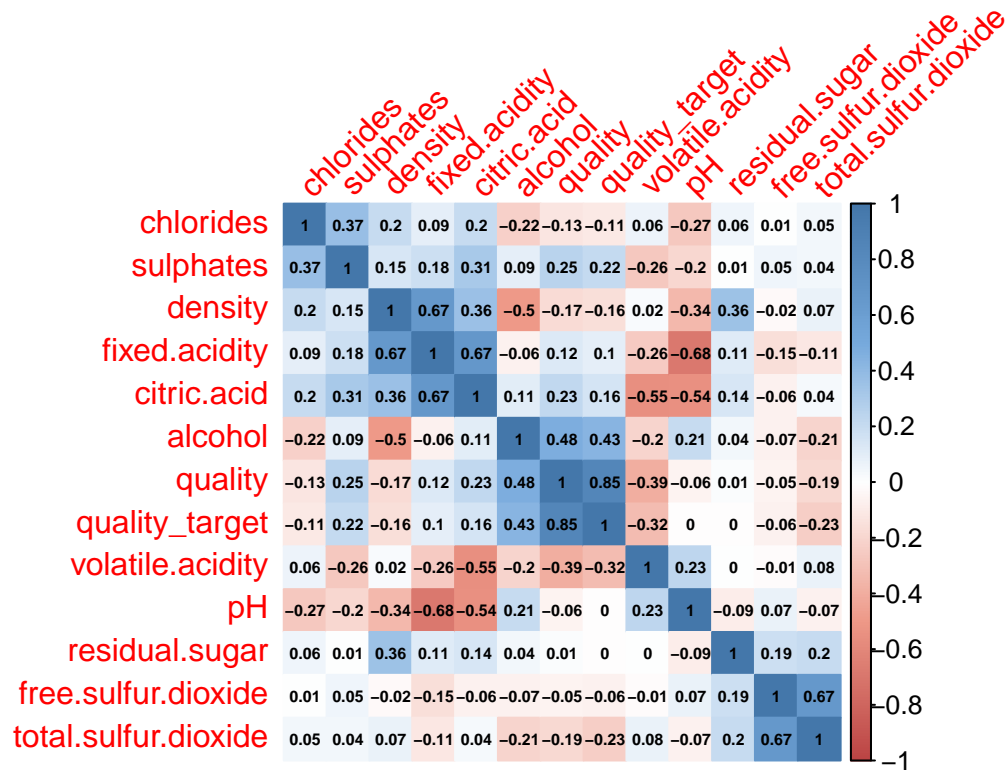
```
# Correlation Matrix
```

```
cor <- cor(wine)
```

```
# Colors for Correlation Matrix
```

```
colors <- colorRampPalette(c("#BB4444", "#EE9988", "#FFFFFF", "#77AADD", "#4477AA"))
```

```
corrplot(cor, order="hclust", method = "color", addCoef.col = "black",  
          , tl.srt = 45, number.cex = 0.47, col=colors(200))
```



```
# Cutoff Correlation features
cutoffCorr <- findCorrelation(cor, cutoff = .8)
cutoffCorrFeatures <- wine[, -cutoffCorr]


# Train and Test split
wine_split <- createDataPartition(wine$quality, p = .8, list = FALSE)
wine_train <- wine[ wine_split,]
wine_test  <- wine[-wine_split,]

# Transform Train Data
train_trans <- preProcess(wine_train, method = c("center", "scale"))
train_transformed <- predict(train_trans, wine_train)

# Transform Test Data
test_trans <- preProcess(wine_test, method = c("center", "scale"))
test_transformed <- predict(test_trans, wine_test)

# Boxplot of transformed train data
boxplot(train_transformed, horizontal = TRUE, las = 2, cex.axis = .65, cex.lab = 7)
```





## Logistic Regression Model

*# Cutoff Correlation string to copy + paste into feature area of model*

```
subset(cutoffCorrFeatures, select = -c(quality_target)) %>%
  colnames() %>%
  paste0(collapse = " + ")
```

```
## [1] "fixed.acidity + volatile.acidity + citric.acid + residual.sugar + chlorides + free.sulfur.dioxide"
set.seed(4)
```

*# Model using "quality\_target" as target variable*

```
lmodel1 <- lm(quality_target ~ volatile.acidity + sulphates + alcohol, data = train_transformed)
```

```
summary(lmodel1)
```

```
##
```

```
## Call:
```

```
## lm(formula = quality_target ~ volatile.acidity + sulphates +
##     alcohol, data = train_transformed)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -2.07422 -0.69261 -0.05795  0.76130  2.13602
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -3.818e-15  2.385e-02   0.000      1
## volatile.acidity -2.337e-01  2.509e-02 -9.313 < 2e-16 ***
## sulphates       1.220e-01  2.470e-02  4.938 8.93e-07 ***
```

```

## alcohol          3.841e-01  2.431e-02  15.796  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8535 on 1277 degrees of freedom
## Multiple R-squared:  0.2733, Adjusted R-squared:  0.2716
## F-statistic: 160.1 on 3 and 1277 DF,  p-value: < 2.2e-16

# Model using "quality" as target variable
lmodel2 <- lm(quality~ volatile.acidity + sulphates + alcohol, data = train_transformed)

summary(lmodel2)

##
## Call:
## lm(formula = quality ~ volatile.acidity + sulphates + alcohol,
##     data = train_transformed)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.01860 -0.47772 -0.06938  0.56533  2.72662
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.850e-15  2.261e-02   0.000      1
## volatile.acidity -2.885e-01  2.379e-02 -12.129 < 2e-16 ***
## sulphates       1.371e-01  2.342e-02   5.856 6.03e-09 ***
## alcohol         4.124e-01  2.305e-02  17.892 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8091 on 1277 degrees of freedom
## Multiple R-squared:  0.3469, Adjusted R-squared:  0.3453
## F-statistic: 226.1 on 3 and 1277 DF,  p-value: < 2.2e-16


# Add predicted values to new data frame
wine_test %>%
  mutate(predicted = predict(lmodel2, newdata = test_transformed)) -> df

# Summary of predicted interval
predict(lmodel2, newdata = test_transformed, interval = "prediction") %>%
  summary()

##           fit           lwr           upr
## Min.      :-1.5138   Min.    :-3.1081   Min.     :0.08058
## 1st Qu.: -0.4639   1st Qu.: -2.0526   1st Qu.: 1.12491
## Median:  -0.1398   Median:  -1.7291   Median:  1.44960
## Mean:     0.0000   Mean:     -1.5898   Mean:     1.58979
## 3rd Qu.:  0.4498   3rd Qu.: -1.1390   3rd Qu.: 2.03853
## Max.      : 2.1802   Max.       0.5808   Max.      :3.77952

# Scatter plot of predicted
ggplot(df, aes(x = predicted, y = quality, colour = quality)) +
  geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()

```



*# The scatter plot supports the summary of the predicted interval, in the ranges of the fit, lower, and upper ranges. The R-squared value of 0.3283 of the model, indicates that this information can be predicted 33% of the time, with the data available, for the variance of the information.*

## CART

```
set.seed(4)
# Subset both train and test sets, to exclude "quality_target"
subset(train_transformed, select = -c(quality_target)) -> rf_wine_train
subset(test_transformed, select = -c(quality_target)) -> rf_wine_test



rPartTree <- rpart(quality ~ ., data = rf_wine_train)

rpartTree2 <- as.party(rPartTree)

# R-Squared plot
par(mfrow=c(1,2))
rsq.rpart(rPartTree)

##
## Regression tree:
## rpart(formula = quality ~ ., data = rf_wine_train)
##
## Variables actually used in tree construction:
```

```
## [1] alcohol          sulphates          total.sulfur.dioxide
## [4] volatile.acidity
##
## Root node error: 1280/1281 = 0.99922
##
## n= 1281
##
##      CP nsplit rel error  xerror  xstd
## 1  0.185022    0  1.00000  1.00121  0.041490
## 2  0.055480    1  0.81498  0.83492  0.040396
## 3  0.033497    2  0.75950  0.80147  0.035118
## 4  0.027105    3  0.72600  0.77633  0.034420
## 5  0.019523    4  0.69890  0.75879  0.033713
## 6  0.019476    5  0.67937  0.75152  0.033808
## 7  0.017209    6  0.65990  0.74874  0.033799
## 8  0.011818    7  0.64269  0.72426  0.031918
## 9  0.011508    8  0.63087  0.71478  0.031714
## 10 0.010997    9  0.61936  0.71066  0.031343
## 11 0.010000   10  0.60837  0.70556  0.030991
```



```
# Results
rpartTree2
```

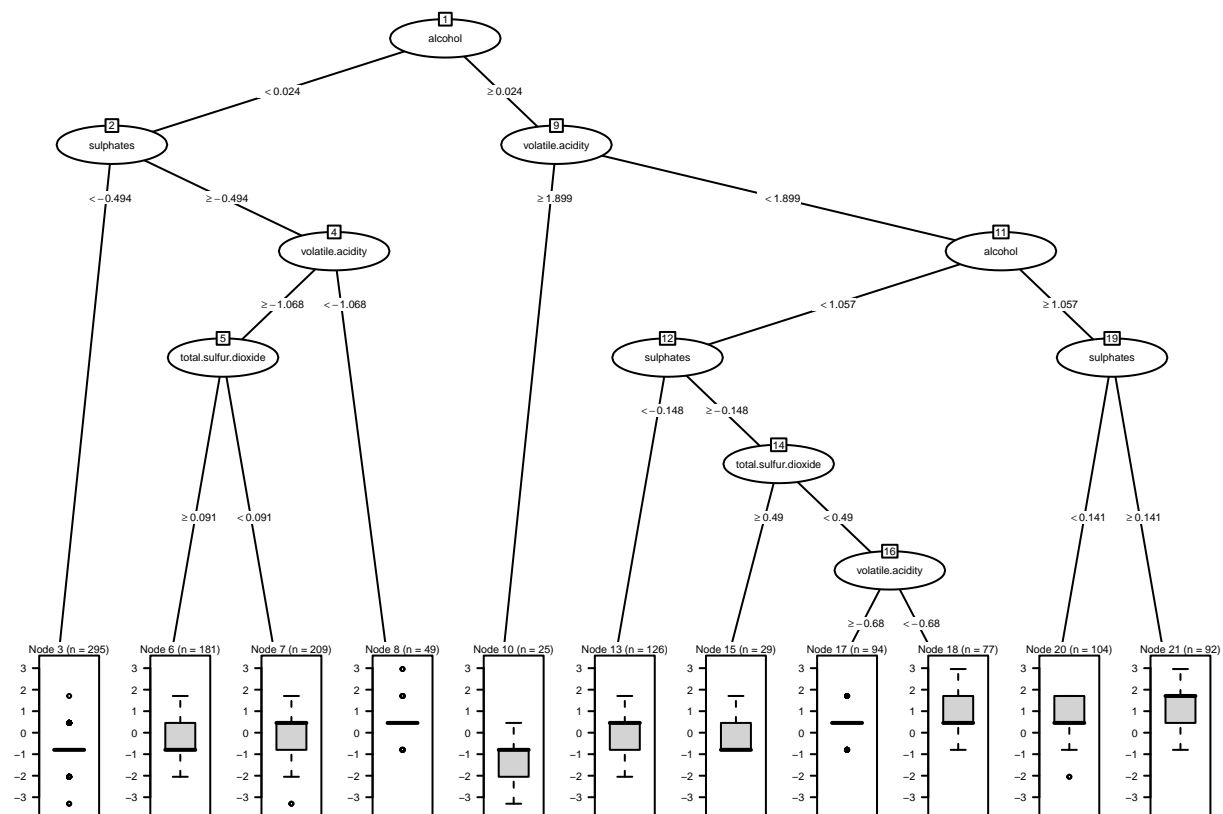
```
##
## Model formula:
## quality ~ fixed.acidity + volatile.acidity + citric.acid + residual.sugar +
## chlorides + free.sulfur.dioxide + total.sulfur.dioxide +
## density + pH + sulphates + alcohol
##
## Fitted party:
## [1] root
## | [2] alcohol < 0.02403
## | | [3] sulphates < -0.49395: -0.636 (n = 295, err = 133.5)
```

```

## |   |   [4] sulphates >= -0.49395
## |   |   [5] volatile.acidity >= -1.06828
## |   |   [6] total.sulfur.dioxide >= 0.09106: -0.486 (n = 181, err = 81.3)
## |   |   [7] total.sulfur.dioxide < 0.09106: -0.097 (n = 209, err = 140.4)
## |   |   [8] volatile.acidity < -1.06828: 0.480 (n = 49, err = 36.0)
## | [9] alcohol >= 0.02403
## |   |   [10] volatile.acidity >= 1.89857: -1.148 (n = 25, err = 33.0)
## |   |   [11] volatile.acidity < 1.89857
## |   |   [12] alcohol < 1.05677
## |   |   [13] sulphates < -0.14788: 0.027 (n = 126, err = 97.7)
## |   |   [14] sulphates >= -0.14788
## |   |   [15] total.sulfur.dioxide >= 0.49048: -0.107 (n = 29, err = 20.7)
## |   |   [16] total.sulfur.dioxide < 0.49048
## |   |   [17] volatile.acidity >= -0.6801: 0.415 (n = 94, err = 48.5)
## |   |   [18] volatile.acidity < -0.6801: 0.991 (n = 77, err = 54.7)
## |   |   [19] alcohol >= 1.05677
## |   |   [20] sulphates < 0.1405: 0.611 (n = 104, err = 74.3)
## |   |   [21] sulphates >= 0.1405: 1.326 (n = 92, err = 58.8)
##
## Number of inner nodes: 10
## Number of terminal nodes: 11

```

```
plot(rpartTree2, gp = gpar(fontsize=4))
```



## Random Forest

```
set.seed(4)

rf <- rfsrc(quality ~ ., data = rf_wine_train)

print(rf)

##                      Sample size: 1281
##                      Number of trees: 500
##                      Forest terminal node size: 5
##                      Average no. of terminal nodes: 152.262
## No. of variables tried at each split: 4
##                      Total no. of variables: 11
##                      Resampling used to grow trees: swor
##                      Resample size used to grow trees: 810
##                      Analysis: RF-R
##                      Family: regr
##                      Splitting rule: mse
##                      (OOB) R squared: 0.47760991
##                      (OOB) Requested performance error: 0.52239009
```

### *# Variable Importance*


```
vi <- subsample(rf, verbose = FALSE)
```

```
extract.subsample(vi)$var.jk.sel.Z
```

	lower	mean	upper	pvalue	signif
## fixed.acidity	34.98371	47.92706	60.87042	1.972777e-13	TRUE
## volatile.acidity	84.51964	110.98665	137.45365	1.026847e-16	TRUE
## citric.acid	17.56287	32.17938	46.79589	7.979662e-06	TRUE
## residual.sugar	41.31289	58.83278	76.35266	2.325943e-11	TRUE
## chlorides	33.83073	49.88923	65.94774	5.678963e-10	TRUE
## free.sulfur.dioxide	16.82409	28.16609	39.50809	5.657645e-07	TRUE
## total.sulfur.dioxide	39.20016	51.85320	64.50624	4.791123e-16	TRUE
## density	34.23564	44.25832	54.28099	2.467949e-18	TRUE
## pH	31.91963	45.66939	59.41915	3.758552e-11	TRUE
## sulphates	45.78972	60.60198	75.41424	5.335758e-16	TRUE
## alcohol	127.23911	166.59546	205.95181	5.359853e-17	TRUE

### *# Variable Importance Plot*

```
plot(vi)
```



```
# Predict
# https://www.rdocumentation.org/packages/randomForestSRC/versions/3.1.0/topics/predict.rfsrc
randomForestSRC::predict.rfsrc(rf, rf_wine_test)
```

```
## Sample size of test (predict) data: 318
## Number of grow trees: 500
## Average no. of grow terminal nodes: 152.262
## Total no. of grow variables: 11
## Resampling used to grow trees: swor
## Resample size used to grow trees: 810
## Analysis: RF-R
## Family: regr
## R squared: 0.41116769
## Requested performance error: 0.58883231
```

## Partial Least Squares

```
tctrl <- trainControl(method = "repeatedcv", repeats = 5, number = 10)

set.seed(4)
pls_wine <- train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data = train_transformed,
  method = "pls",
  preProc = c("center", "scale", "BoxCox"),
  tunelength = 20,
  trControl = tctrl)

pls_wine
```

```
## Partial Least Squares
##
```

```
## 1281 samples
## 5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 1153, 1153, 1153, 1154, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      0.7998011  0.3632430  0.6276163
##   2      0.7993504  0.3641251  0.6271541
##   3      0.7994295  0.3640622  0.6267032
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 2.
```

## Mars Tuning

```
mars_wine <- earth(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data=train_transformed)
```

```
mars_wine
```

```
## Selected 15 of 16 terms, and 5 of 5 predictors
## Termination condition: Reached nk 21
## Importance: alcohol, volatile.acidity, sulphates, total.sulfur.dioxide, ...
## Number of terms at each degree of interaction: 1 14 (additive model)
## GCV 0.6128238   RSS 749.8859   GRSq 0.3876546   RSq 0.4141517
```

```
summary(mars_wine)
```

```
## Call: earth(formula=quality~volatile.acidity+chlorides+total.sulfur.di...),
## data=train_transformed)
```

```
##
## coefficients
## (Intercept) 30.4984030
## h(1.73221-volatile.acidity) 0.1976638
## h(volatile.acidity-1.73221) -0.4644967
## h(chlorides- -0.960266) 3.2533522
## h(chlorides-0.0972729) -0.6250823
## h(1.39886-chlorides) 3.0770129
## h(chlorides-1.39886) -2.4570771
## h(chlorides-5.14092) -0.4076001
## h(total.sulfur.dioxide- -1.1368) -10.2398712
## h(2.53198-total.sulfur.dioxide) -10.0963829
## h(total.sulfur.dioxide-2.53198) 10.5813790
## h(1.0345-sulphates) -0.4128928
## h(sulphates-1.0345) -0.1156520
## h(alcohol-0.634282) 0.2129655
## h(1.76091-alcohol) -0.3017427
##
```

```
## Selected 15 of 16 terms, and 5 of 5 predictors
## Termination condition: Reached nk 21
```



```

## Importance: alcohol, volatile.acidity, sulphates, total.sulfur.dioxide, ...
## Number of terms at each degree of interaction: 1 14 (additive model)
## GCV 0.6128238    RSS 749.8859    GRSq 0.3876546    RSq 0.4141517

preProc_Arguments = c("center", "scale")
marsGrid_wine = expand.grid(.degree=1:2, .nprune=2:38)

set.seed(4)

marsModel_wine = train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data =train_transformed,
  method="earth",
  preProc=preProc_Arguments,
  tuneGrid=marsGrid_wine)

marsModel_wine

## Multivariate Adaptive Regression Spline
##
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1281, 1281, 1281, 1281, 1281, 1281, ...
## Resampling results across tuning parameters:
##
##  degree  nprune  RMSE      Rsquared  MAE
##  1         2     0.8767725  0.2286331  0.6983669
##  1         3     0.8281246  0.3129505  0.6469912
##  1         4     0.7986723  0.3598732  0.6239193
##  1         5     0.7973767  0.3627609  0.6213442
##  1         6     0.7954260  0.3658170  0.6189612
##  1         7     0.7933841  0.3690469  0.6184932
##  1         8     0.7955313  0.3658275  0.6195302
##  1         9     0.7933719  0.3687607  0.6178130
##  1        10     0.7943316  0.3680007  0.6182851
##  1        11     0.7949170  0.3675983  0.6183133
##  1        12     0.7955150  0.3670106  0.6189397
##  1        13     0.7956669  0.3668727  0.6185138
##  1        14     0.7961519  0.3667770  0.6185384
##  1        15     0.7962221  0.3667518  0.6180695
##  1        16     0.7962221  0.3667518  0.6180695
##  1        17     0.7962221  0.3667518  0.6180695
##  1        18     0.7962221  0.3667518  0.6180695
##  1        19     0.7962221  0.3667518  0.6180695
##  1        20     0.7962221  0.3667518  0.6180695
##  1        21     0.7962221  0.3667518  0.6180695
##  1        22     0.7962221  0.3667518  0.6180695
##  1        23     0.7962221  0.3667518  0.6180695
##  1        24     0.7962221  0.3667518  0.6180695
##  1        25     0.7962221  0.3667518  0.6180695
##  1        26     0.7962221  0.3667518  0.6180695
##  1        27     0.7962221  0.3667518  0.6180695
##  1        28     0.7962221  0.3667518  0.6180695

```



## KNN Neighbors

```
set.seed(4)

knn_wine <- train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data =train_transformed,
  method = "knn",
  preProc = c("center", "scale"),
  tuneGrid = data.frame(.k = 1:50),
  trControl = trainControl(method = "cv"))

knn_wine
```

```
## k-Nearest Neighbors
##
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1153, 1153, 1153, 1154, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##   k  RMSE      Rsquared  MAE
##   1  0.9486037  0.3101279  0.5521179
##   2  0.8946471  0.3055655  0.6207649
##   3  0.8596248  0.3173637  0.6192937
##   4  0.8384012  0.3314930  0.6215991
##   5  0.8202779  0.3497374  0.6164375
##   6  0.8096196  0.3604723  0.6099141
##   7  0.7961261  0.3766117  0.6026718
##   8  0.7917952  0.3808891  0.6029282
##   9  0.7930582  0.3770572  0.6082138
##  10  0.7915187  0.3787846  0.6105603
##  11  0.7868061  0.3848931  0.6104114
##  12  0.7883179  0.3813246  0.6116129
##  13  0.7877965  0.3815857  0.6121286
##  14  0.7869724  0.3826114  0.6130522
##  15  0.7867557  0.3823977  0.6126466
##  16  0.7862693  0.3831196  0.6100775
##  17  0.7877585  0.3804738  0.6100929
##  18  0.7882592  0.3799676  0.6112571
##  19  0.7858700  0.3836689  0.6103664
##  20  0.7853813  0.3844932  0.6113662
##  21  0.7863263  0.3826412  0.6120682
##  22  0.7866219  0.3821987  0.6121664
##  23  0.7864929  0.3826019  0.6124847
##  24  0.7866213  0.3827113  0.6129548
##  25  0.7839553  0.3868342  0.6113461
##  26  0.7830884  0.3882729  0.6103765
##  27  0.7840037  0.3868623  0.6121240
##  28  0.7832355  0.3877893  0.6114887
##  29  0.7828527  0.3886391  0.6111209
##  30  0.7815061  0.3908631  0.6108329
```

```
## 31 0.7816921 0.3905510 0.6112563
## 32 0.7813698 0.3911467 0.6116271
## 33 0.7824046 0.3896750 0.6119544
## 34 0.7812635 0.3915225 0.6107396
## 35 0.7821713 0.3899855 0.6115484
## 36 0.7829359 0.3888554 0.6126616
## 37 0.7826544 0.3894240 0.6128846
## 38 0.7825614 0.3898820 0.6133241
## 39 0.7824786 0.3901287 0.6128711
## 40 0.7819865 0.3911644 0.6130563
## 41 0.7831124 0.3894069 0.6142938
## 42 0.7831814 0.3892890 0.6141271
## 43 0.7825908 0.3904415 0.6145867
## 44 0.7828766 0.3900157 0.6150421
## 45 0.7832661 0.3894957 0.6155824
## 46 0.7831180 0.3896859 0.6153044
## 47 0.7826061 0.3905436 0.6153951
## 48 0.7814354 0.3926554 0.6139669
## 49 0.7817824 0.3922137 0.6144282
## 50 0.7819978 0.3919336 0.6143947
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 34.
```

## SVM

```
set.seed(4)

subset(train_transformed, select = -c(quality_target, quality)) -> predictors
train_transformed$quality -> quality

svmTune <- train(predictors, quality,
                 method = "svmRadial",
                 preProc = c("center", "scale"),
                 tuneLength= 5,
                 trControl = trainControl(method = "cv"))
svmTune

## Support Vector Machines with Radial Basis Function Kernel
##
## 1281 samples
## 11 predictor
##
## Pre-processing: centered (11), scaled (11)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1153, 1153, 1153, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##  C      RMSE      Rsquared  MAE
##  0.25  0.7807204  0.3936554  0.5841701
##  0.50  0.7724359  0.4053327  0.5722882
##  1.00  0.7748364  0.4034222  0.5722157
```

```
## 2.00 0.7804447 0.3991023 0.5747985
## 4.00 0.7910344 0.3920209 0.5818857
##
## Tuning parameter 'sigma' was held constant at a value of 0.1020787
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.1020787 and C = 0.5.
```