

ADS 503 - Team 7

Summer Purschke, Jacqueline Urenda, Oscar Gil

06/12/2022

```
# R Libraries
library(caret)
library(AppliedPredictiveModeling)
library(Hmisc)
library(dplyr)
library(tidyverse)
library(ggplot2)
library(corrplot)
library(MASS)
library(ISLR)
library(rpart)
library(partykit)
library(randomForestSRC)
library(earth)
library(MARSS)
library(e1071)
library(summarytools)
library(grid)
library(MLeval)
library(pROC)
```

Load the Red Wine Quality data set from GitHub - data set copied from Kaggle and imported into GitHub.

```
wine <- read.csv(
  url("https://raw.githubusercontent.com/OscarG-DataSci/ADS503/main/winequality-red.csv")
  , header = TRUE)
```

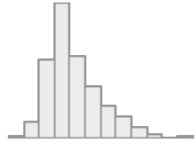
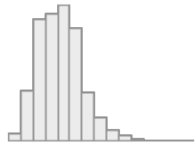
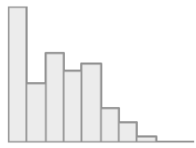
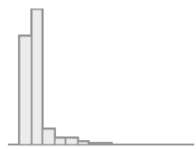
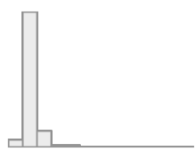
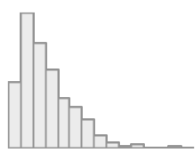
Data Summary

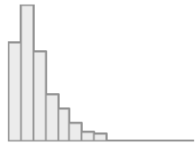
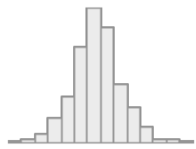
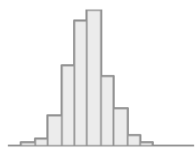
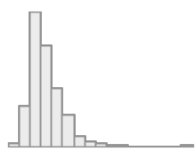
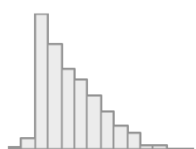
```
# use the view function to view in R Studio
#view(
dfSummary(wine,
  plain.ascii = FALSE,
  style       = "grid",
  graph.magnif = 0.75,
  valid.col   = FALSE,
  tmp.img.dir = "NA")
```

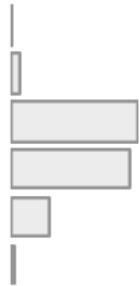
Data Frame Summary

wine Dimensions: 1599 x 12

Duplicates: 240

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
1	fixed.acidity [numeric]	Mean (sd) : 8.3 (1.7) min < med < max: 4.6 < 7.9 < 15.9 IQR (CV) : 2.1 (0.2)	96 distinct values		0 (0.0%)
2	volatile.acidity [numeric]	Mean (sd) : 0.5 (0.2) min < med < max: 0.1 < 0.5 < 1.6 IQR (CV) : 0.2 (0.3)	143 distinct values		0 (0.0%)
3	citric.acid [numeric]	Mean (sd) : 0.3 (0.2) min < med < max: 0 < 0.3 < 1 IQR (CV) : 0.3 (0.7)	80 distinct values		0 (0.0%)
4	residual.sugar [numeric]	Mean (sd) : 2.5 (1.4) min < med < max: 0.9 < 2.2 < 15.5 IQR (CV) : 0.7 (0.6)	91 distinct values		0 (0.0%)
5	chlorides [numeric]	Mean (sd) : 0.1 (0) min < med < max: 0 < 0.1 < 0.6 IQR (CV) : 0 (0.5)	153 distinct values		0 (0.0%)
6	free.sulfur.dioxide [numeric]	Mean (sd) : 15.9 (10.5) min < med < max: 1 < 14 < 72 IQR (CV) : 14 (0.7)	60 distinct values		0 (0.0%)

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
7	total.sulfur.dioxide [numeric]	Mean (sd) : 46.5 (32.9) min < med < max: 6 < 38 < 289 IQR (CV) : 40 (0.7)	144 distinct values		0 (0.0%)
8	density [numeric]	Mean (sd) : 1 (0) min < med < max: 1 < 1 < 1 IQR (CV) : 0 (0)	436 distinct values		0 (0.0%)
9	pH [numeric]	Mean (sd) : 3.3 (0.2) min < med < max: 2.7 < 3.3 < 4 IQR (CV) : 0.2 (0)	89 distinct values		0 (0.0%)
10	sulphates [numeric]	Mean (sd) : 0.7 (0.2) min < med < max: 0.3 < 0.6 < 2 IQR (CV) : 0.2 (0.3)	96 distinct values		0 (0.0%)
11	alcohol [numeric]	Mean (sd) : 10.4 (1.1) min < med < max: 8.4 < 10.2 < 14.9 IQR (CV) : 1.6 (0.1)	65 distinct values		0 (0.0%)

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
					
12	quality [integer]	Mean (sd) : 5.6 (0.8) min < med < max: 3 < 6 < 8 IQR (CV) : 1 (0.1)	3 : 10 (0.6%) 4 : 53 (3.3%) 5 : 681 (42.6%) 6 : 638 (39.9%) 7 : 199 (12.4%) 8 : 18 (1.1%)		0 (0.0%)

```
# )
```

Pre-processing

```
par(mar=c(1,1,1,1)) # to fix boxplot knit processing issues
```

```
# Create new variable, for quality values, split by half (0, 1)
wine$quality_target <- ifelse( wine$quality <= 5, 0, 1)
```

```
# Mean of new variable is at 0.5347 (close enough to 50% to maintain balance)
summary(wine$quality_target)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000 0.0000  1.0000  0.5347  1.0000  1.0000
```

```
# Check for missing values in data set
wine %>% na.omit() %>% count() # there are no missing values
```

```
##      n
## 1 1599
```

```
# Removing outliers for residual sugar:
```

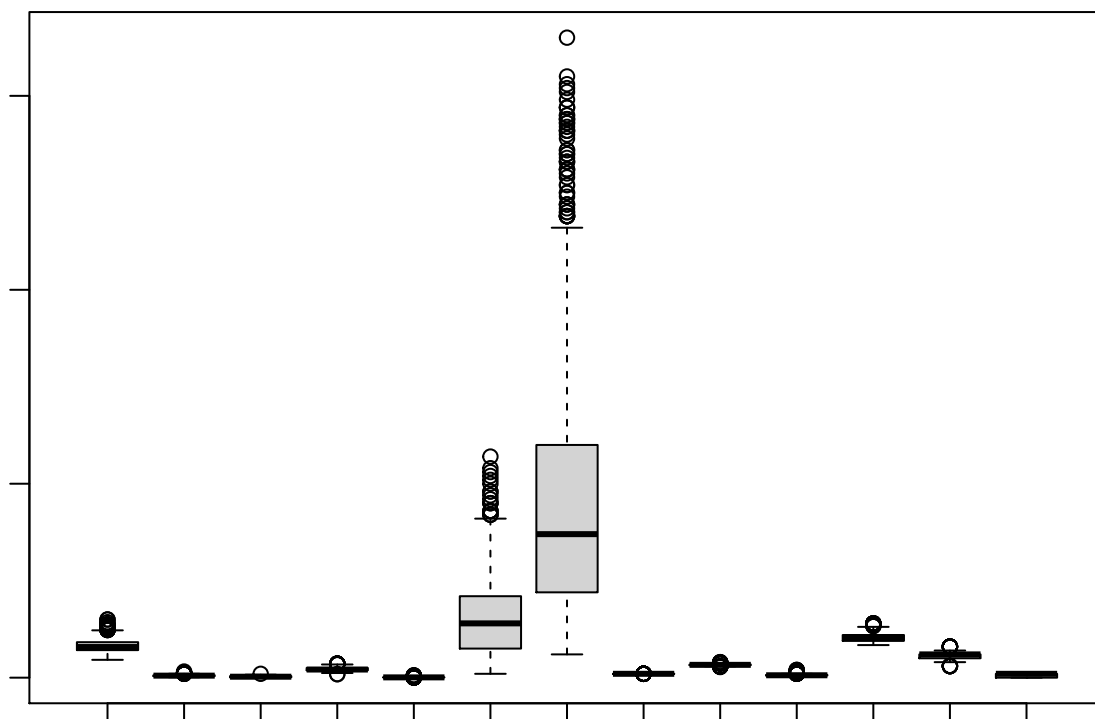
```
Q <- quantile(wine$residual.sugar, probs=c(.25, .75), na.rm = FALSE)
```

```
iqr_rs <- IQR(wine$residual.sugar)
```

```
up_rs <- Q[2]+1.5*iqr_rs # Upper Range
```

```
low_rs <- Q[1]-1.5*iqr_rs # Lower Range
```

```
eliminated_rs <- subset(wine, wine$residual.sugar > (Q[1] - 1.5*iqr_rs) & wine$residual.sugar < (Q[2]+1.5*iqr_rs))
boxplot(eliminated_rs)
```



#Removing outliers for free.sulfur.dioxide:

```
Q2 <- quantile(wine$free.sulfur.dioxide, probs=c(.25, .75), na.rm = FALSE)
```

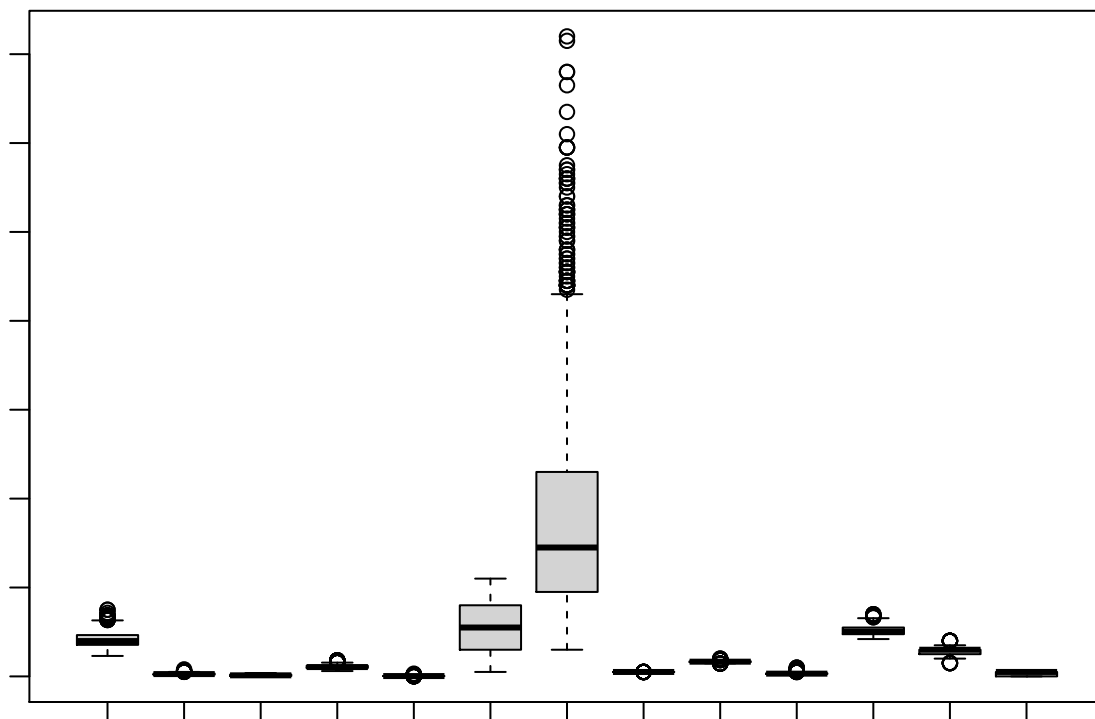
```
iqr_fs <- IQR(eliminated_rs$free.sulfur.dioxide)
```

```
up_fs <- Q2[2]+1.5*iqr_fs # Upper Range
```

```
low_fs <- Q2[1]-1.5*iqr_fs # Lower Range
```

```
eliminated_fs <- subset(eliminated_rs, eliminated_rs$free.sulfur.dioxide > (Q[1] - 1.5*iqr_fs) & elimin
```

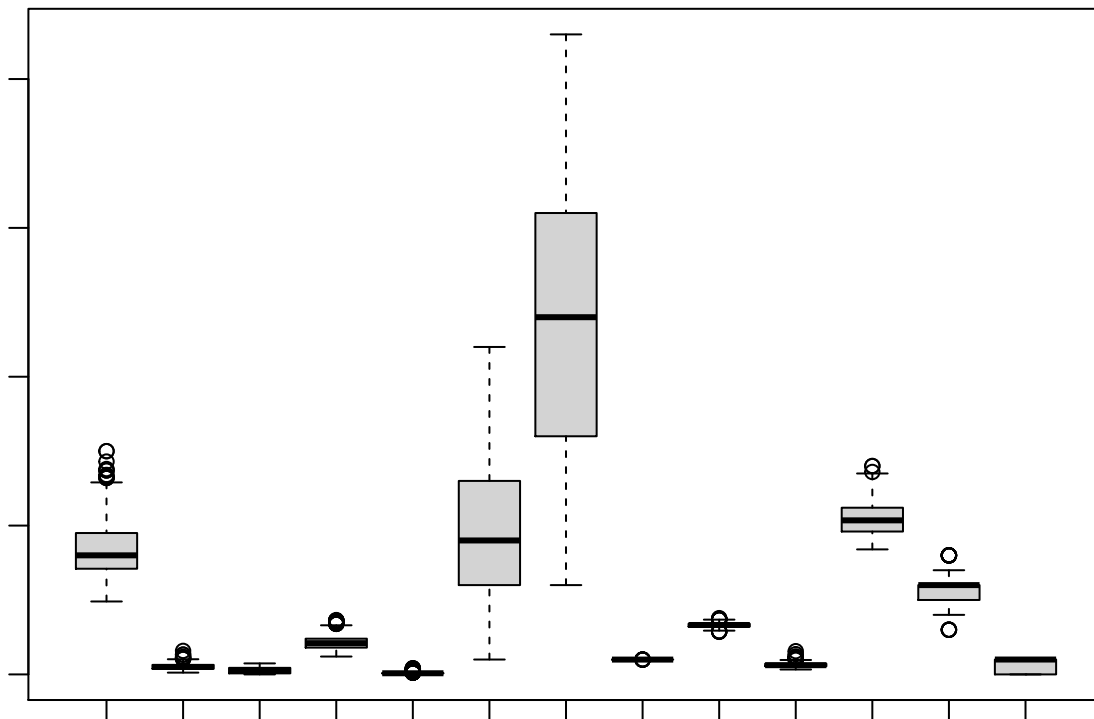
```
boxplot(eliminated_fs)
```



```

#Removing outliers for total.sulfur.dioxide:
Q3 <- quantile(wine$total.sulfur.dioxide, probs=c(.25, .75), na.rm = FALSE)
iqr_ts <- IQR(eliminated_fs$total.sulfur.dioxide)
up_ts <- Q3[2]+1.5*iqr_ts # Upper Range
low_ts <- Q3[1]-1.5*iqr_ts # Lower Range
eliminated_ts <- subset(eliminated_fs, eliminated_fs$total.sulfur.dioxide > (Q[1] - 1.5*iqr_ts) & eliminated_ts < up_ts)
boxplot(eliminated_ts)

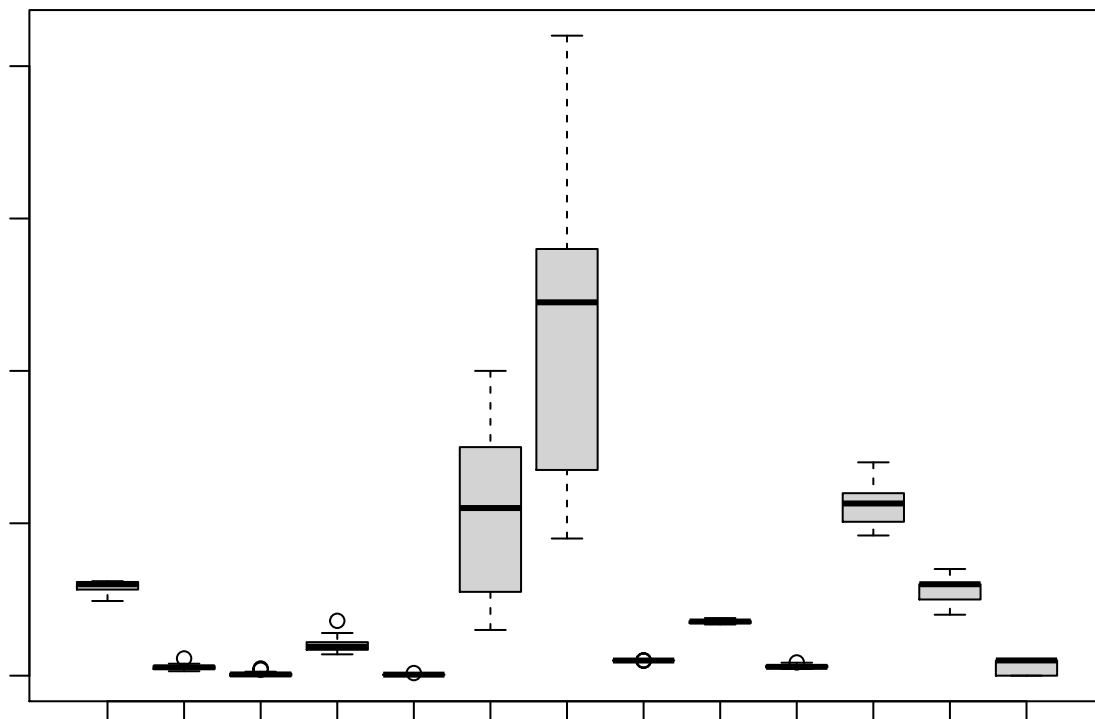
```



```

#Removing outliers for fixed.acidity:
Q4 <- quantile(wine$fixed.acidity, probs=c(.25, .75), na.rm = FALSE)
iqr_fa <- IQR(eliminated_ts$fixed.acidity)
up_fa <- Q[2]+1.5*iqr_fa # Upper Range
low_fa <- Q[1]-1.5*iqr_fa # Lower Range
eliminated_fa <- subset(eliminated_ts, eliminated_ts$fixed.acidity > (Q[1] - 1.5*iqr_fa) & eliminated_ts < up_fa)
boxplot(eliminated_fa)

```



```
new_wine_data <- eliminated_fa
```

```
# Removing outliers reduced dimension of data set from 1599 observations to 48
```

```
# team opted not to use new_wine_data and keep outlier data
```

```
dim(new_wine_data)
```

```
## [1] 48 13
```

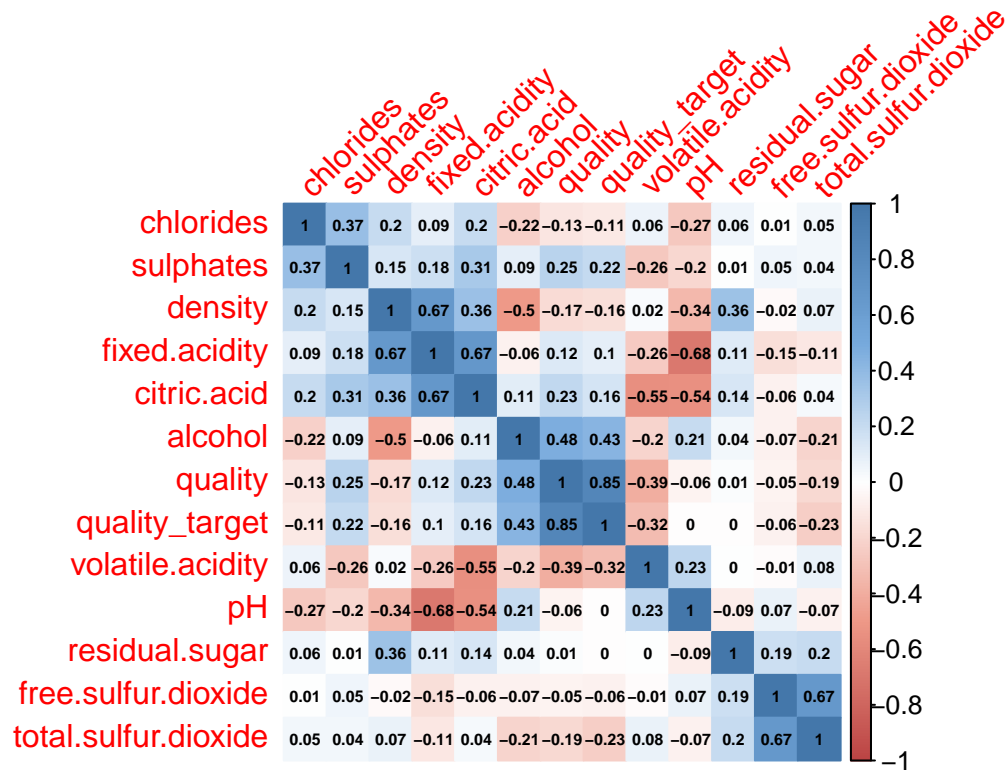
```
# Correlation Matrix
```

```
cor <- cor(wine)
```

```
# Colors for Correlation Matrix
```

```
colors <- colorRampPalette(c("#BB4444", "#EE9988", "#FFFFFF", "#77AADD", "#4477AA"))
```

```
corrplot(cor, order="hclust", method = "color", addCoef.col = "black",  
          , tl.srt = 45, number.cex = 0.47, col=colors(200))
```




```
# Cutoff Correlation features
cutoffCorr <- findCorrelation(cor, cutoff = .8)
cutoffCorrFeatures <- wine[, -cutoffCorr]

# Train and Test split
wine_split <- createDataPartition(wine$quality, p = .8, list = FALSE)
wine_train <- wine[ wine_split,]
wine_test  <- wine[-wine_split,]

# Transform Train Data
train_trans <- preProcess(wine_train, method = c("center", "scale"))
train_transformed <- predict(train_trans, wine_train)

# Transform Test Data
test_trans <- preProcess(wine_test, method = c("center", "scale"))
test_transformed <- predict(test_trans, wine_test)

# Boxplot of transformed train data
boxplot(train_transformed, horizontal = TRUE, las = 2, cex.axis = .65, cex.lab = 7)
```

Logistic Regression Model

```
# Cutoff Correlation string to copy + paste into feature area of model
subset(cutoffCorrFeatures, select = -c(quality_target)) %>%
  colnames() %>%
  paste0(collapse = " + ")
```

```
## [1] "fixed.acidity + volatile.acidity + citric.acid + residual.sugar + chlorides + free.sulfur.dioxide"
set.seed(4)
```

```
# Model using "quality_target" as target variable
lmodel1 <- lm(quality_target ~ volatile.acidity + sulphates + alcohol, data = wine_train)

summary(lmodel1)
```

```
##
## Call:
## lm(formula = quality_target ~ volatile.acidity + sulphates +
##     alcohol, data = wine_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.49365 -0.35478 -0.02426  0.38227  1.03461
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -1.23371    0.14306  -8.624  < 2e-16 ***
## volatile.acidity -0.59164    0.07143  -8.283 3.00e-16 ***
## sulphates       0.34458    0.07124   4.837 1.48e-06 ***
```

```

## alcohol          0.17791    0.01154  15.424 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4292 on 1277 degrees of freedom
## Multiple R-squared:  0.262, Adjusted R-squared:  0.2603
## F-statistic: 151.1 on 3 and 1277 DF,  p-value: < 2.2e-16

# Model using "quality" as target variable
lmodel2 <- lm(quality~ volatile.acidity + sulphates + alcohol, data = wine_train)

summary(lmodel2)

##
## Call:
## lm(formula = quality ~ volatile.acidity + sulphates + alcohol,
##     data = wine_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7039 -0.3831 -0.0719  0.4776  2.1995
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.62659    0.21803   12.047 < 2e-16 ***
## volatile.acidity -1.16761    0.10886  -10.726 < 2e-16 ***
## sulphates       0.60363    0.10858    5.559 3.29e-08 ***
## alcohol         0.31030    0.01758   17.651 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.6541 on 1277 degrees of freedom
## Multiple R-squared:  0.3328, Adjusted R-squared:  0.3313
## F-statistic: 212.4 on 3 and 1277 DF,  p-value: < 2.2e-16

# Add predicted values to new data frame
wine_test %>%
  mutate(predicted = predict(lmodel2, newdata = wine_test)) -> df

# Summary of predicted interval
predict(lmodel2, newdata = wine_test, interval = "prediction") %>%
  summary()

##           fit           lwr           upr
## Min.       :4.405   Min.     :3.103   Min.     :5.708
## 1st Qu.:5.304   1st Qu.:4.019   1st Qu.:6.588
## Median :5.628   Median :4.343   Median :6.914
## Mean    :5.645   Mean    :4.360   Mean     :6.930
## 3rd Qu.:5.957   3rd Qu.:4.672   3rd Qu.:7.241
## Max.    :7.051   Max.    :5.761   Max.     :8.340

# Confusion Matrix
# Convert predicted values to whole numbers, so they match target values
df$predicted_int = as.integer(round(df$predicted, digits = 0))

union1 <- union(df$quality, df$predicted_int)

```

```

table1 <- table(factor(df$quality, union1), factor(df$predicted_int, union1))


confusionMatrix(table1)

## Confusion Matrix and Statistics
##
##
##      5  6  4  7  8  3
## 5 87 46  0  0  0  0
## 6 39 83  0  5  0  0
## 4  8  2  0  0  0  0
## 7  1 33  0  6  0  0
## 8  0  3  0  0  0  0
## 3  4  0  1  0  0  0
##
## Overall Statistics
##
##              Accuracy : 0.5535
##              95% CI : (0.497, 0.6089)
##      No Information Rate : 0.5252
##      P-Value [Acc > NIR] : 0.1699
##
##              Kappa : 0.2595
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: 5 Class: 6 Class: 4 Class: 7 Class: 8 Class: 3
## Sensitivity      0.6259   0.4970 0.000000  0.54545      NA      NA
## Specificity      0.7430   0.7086 0.968454  0.88925 0.990566  0.98428
## Pos Pred Value   0.6541   0.6535 0.000000  0.15000      NA      NA
## Neg Pred Value   0.7189   0.5602 0.996753  0.98201      NA      NA
## Prevalence       0.4371   0.5252 0.003145  0.03459 0.000000  0.00000
## Detection Rate   0.2736   0.2610 0.000000  0.01887 0.000000  0.00000
## Detection Prevalence 0.4182  0.3994 0.031447  0.12579 0.009434  0.01572
## Balanced Accuracy 0.6845   0.6028 0.484227  0.71735      NA      NA

# ROC plot
df$predicted_int = round(as.numeric(as.character(df$predicted))), digits = 0)

roc(df$quality, df$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)

```




```
##
## Call:
## roc.default(response = df$quality, predictor = df$predicted_int, percent = TRUE, plot = TRUE, legend = TRUE)
##
## Data: df$predicted_int in 5 controls (df$quality 3) < 10 cases (df$quality 4).
## Area under the curve: 68%

#multiclass.roc(df$quality, df$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)

# modelName1 <- 'Logistic Regression'
# roc1 <- multiclass.roc(df$quality, df$predicted_int)
# auc1 <- round(auc(df$quality, df$predicted_int), 4)
#
# ggroc(roc1, colours = 'red', size = 1) +
#   ggtitle(paste0(modelName1, ' - ROC Curve ', '(AUC = ', auc1, ')')) + theme_minimal()

# Scatter plot of predicted
ggplot(df, aes(x = predicted, y = quality, colour = quality)) +
  geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```



The scatter plot supports the summary of the predicted interval, in the ranges of the fit, lower, and upper ranges. The R-squared value of 0.3283 of the model, indicates that this information can be predicted 33% of the time, with the data available, for the variance of the information.

CART

```
set.seed(4)
# Subset both train and test sets, to exclude "quality_target"
# Using non-transformed versions of train and test, to get actual values in the nodes
subset(wine_train, select = -c(quality_target)) -> rf_wine_train
subset(wine_test, select = -c(quality_target)) -> rf_wine_test

# Convert target variable to factor to ensure proper interpretation by model
rf_wine_train$quality <- as.factor(rf_wine_train$quality)



# Begin model...
rpartTree <- rpart(quality ~ ., data = rf_wine_train)

rpartTree2 <- as.party(rpartTree)

# R-Squared plot
par(mfrow=c(1,2))
rsq.rpart(rpartTree)

##
## Classification tree:
## rpart(formula = quality ~ ., data = rf_wine_train)
```

```
##
## Variables actually used in tree construction:
## [1] alcohol          chlorides          density
## [4] sulphates        total.sulfur.dioxide volatile.acidity
##
## Root node error: 733/1281 = 0.57221
##
## n= 1281
##
##      CP nsplit rel error  xerror   xstd
## 1 0.249659      0  1.00000 1.00000 0.024158
## 2 0.011596      1  0.75034 0.76262 0.024216
## 3 0.010914      5  0.69577 0.74352 0.024141
## 4 0.010459      7  0.67394 0.74352 0.024141
## 5 0.010000     10  0.64256 0.73124 0.024087
```




```
# Results
rpartTree2
```

```
##
## Model formula:
## quality ~ fixed.acidity + volatile.acidity + citric.acid + residual.sugar +
##          chlorides + free.sulfur.dioxide + total.sulfur.dioxide +
##          density + pH + sulphates + alcohol
##
## Fitted party:
## [1] root
## |   [2] alcohol < 10.25
## |   |   [3] total.sulfur.dioxide >= 98.5: 5 (n = 84, err = 7.1%)
## |   |   [4] total.sulfur.dioxide < 98.5
## |   |   |   [5] sulphates < 0.555: 5 (n = 204, err = 27.9%)
## |   |   |   [6] sulphates >= 0.555
## |   |   |   |   [7] volatile.acidity >= 0.365: 5 (n = 336, err = 44.3%)
```

```
## |   |   |   |   [8] volatile.acidity < 0.365: 6 (n = 57, err = 36.8%)
## |   [9] alcohol >= 10.25
## |   |   [10] volatile.acidity >= 0.385
## |   |   |   [11] sulphates < 0.585
## |   |   |   |   [12] density >= 0.99548: 5 (n = 62, err = 45.2%)
## |   |   |   |   [13] density < 0.99548: 6 (n = 59, err = 42.4%)
## |   |   |   [14] sulphates >= 0.585: 6 (n = 278, err = 39.6%)
## |   |   [15] volatile.acidity < 0.385
## |   |   |   [16] total.sulfur.dioxide >= 49.5: 6 (n = 48, err = 41.7%)
## |   |   |   [17] total.sulfur.dioxide < 49.5
## |   |   |   |   [18] sulphates < 0.545: 6 (n = 15, err = 6.7%)
## |   |   |   |   [19] sulphates >= 0.545
## |   |   |   |   |   [20] chlorides >= 0.1105: 6 (n = 12, err = 16.7%)
## |   |   |   |   |   [21] chlorides < 0.1105: 7 (n = 126, err = 41.3%)
##
## Number of inner nodes:    10
## Number of terminal nodes: 11
```

```
plot(rpartTree2, gp = gpar(fontsize=4))
```



```
# Add predicted values to new data frame
```

```
wine_test %>%
```

```
  mutate(predicted = predict(rpartTree2, newdata = wine_test)) -> df2
```

```
# Summary of predicted values
```

```
predict(rpartTree2, newdata = wine_test, interval = "prediction") %>%
  summary()
```

```
##    3    4    5    6    7    8
```

```
## 0 0 167 115 36 0
```

```
# Confusion Matrix
```

```
confusionMatrix(table(df2$quality, df2$predicted))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##
```

```
##      3  4  5  6  7  8
```

```
## 3  0  0  4  1  0  0
```

```
## 4  0  0  7  3  0  0
```

```
## 5  0  0 101 26  6  0
```

```
## 6  0  0  53 58 16  0
```

```
## 7  0  0  2 25 13  0
```

```
## 8  0  0  0  2  1  0
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.5409
```

```
##           95% CI : (0.4844, 0.5966)
```

```
## No Information Rate : 0.5252
```

```
## P-Value [Acc > NIR] : 0.3069
```

```
##
```

```
##           Kappa : 0.2615
```

```
##
```

```
## McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
##           Class: 3 Class: 4 Class: 5 Class: 6 Class: 7 Class: 8
```

```
## Sensitivity           NA           NA  0.6048  0.5043  0.36111           NA
```

```
## Specificity           0.98428  0.96855  0.7881  0.6601  0.90426 0.990566
```

```
## Pos Pred Value           NA           NA  0.7594  0.4567  0.32500           NA
```

```
## Neg Pred Value           NA           NA  0.6432  0.7016  0.91727           NA
```

```
## Prevalence             0.00000  0.00000  0.5252  0.3616  0.11321 0.000000
```

```
## Detection Rate          0.00000  0.00000  0.3176  0.1824  0.04088 0.000000
```

```
## Detection Prevalence    0.01572  0.03145  0.4182  0.3994  0.12579 0.009434
```

```
## Balanced Accuracy           NA           NA  0.6964  0.5822  0.63268           NA
```

```
# ROC plot
```

```
df2$predicted_int = round(as.numeric(as.character(df2$predicted)), digits = 0)
```

```
#roc(df2$quality, df2$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)
```

```
#
```


```
modelName2 <- 'CART'
```

```
roc2 <- roc(df2$quality, df2$predicted_int)
```


```
auc2 <- round(auc(df2$quality, df2$predicted_int), 4)
```

```
ggroc(roc2, colours = 'red', size = 1) +
```

```
  ggtitle(paste0(modelName2, ' - ROC Curve ', '(AUC = ', auc2, ')')) + theme_minimal()
```


```
# Scatter plot of predicted  
ggplot(df2, aes(x = predicted, y = quality, colour = quality ))+  
geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```



```
# Root Node Left vs Right, Quality Density Comparisons
grid.newpage()
filter(wine_train, alcohol < 10.525) %>%
  dplyr::select(quality, alcohol) %>%
  ggplot(aes(x = quality)) + geom_density() -> RootNodeLeft

filter(wine_train, alcohol >= 10.525) %>%
  dplyr::select(quality, alcohol) %>%
  ggplot(aes(x = quality)) + geom_density() -> RootNodeRight

grid.draw(rbind(ggplotGrob(RootNodeLeft), ggplotGrob(RootNodeRight), size = "last"))
```



Random Forest

```
set.seed(4)

rf <- rfsrc(quality ~ ., data = rf_wine_train)

print(rf)
```

```
##                               Sample size: 1281
##           Frequency of class labels: 5, 43, 548, 511, 159, 15
##                   Number of trees: 500
##           Forest terminal node size: 1
##           Average no. of terminal nodes: 252.772
## No. of variables tried at each split: 4
##                   Total no. of variables: 11
##           Resampling used to grow trees: swor
##           Resample size used to grow trees: 810
##                               Analysis: RF-C
##                               Family: class
##                               Splitting rule: gini
##           (OOB) Brier score: 0.07037366
##           (OOB) Normalized Brier score: 0.50669034
##                               (OOB) AUC: 0.78033697
##           (OOB) Requested performance error: 0.30523029, 1, 1, 0.20620438, 0.27201566, 0.49056604, 0.86666666
##
```


```
## Confusion matrix:
##
##      predicted
## observed 3 4  5  6  7  8 class.error
##      3 0 0  5  0  0  0      1.0000
##      4 0 0 29 13  1  0      1.0000
##      5 0 0 436 106  6  0      0.2044
##      6 0 3 113 370 24  1      0.2759
##      7 0 0  9  70 80  0      0.4969
##      8 0 0  0  7  6  2      0.8667
##
##      (OOB) Misclassification rate: 0.3067916
```

```
# Variable Importance
vi <- subsample(rf, verbose = FALSE)

extract.subsample(vi)$var.jk.sel.Z
```

	lower	mean	upper	pvalue	signif
fixed.acidity	6.999708	9.965180	12.930653	2.255101e-11	TRUE
volatile.acidity	12.386745	16.756263	21.125781	2.822027e-14	TRUE
citric.acid	5.687040	8.045860	10.404680	1.151672e-11	TRUE
residual.sugar	8.336291	11.979618	15.622945	5.797002e-11	TRUE
chlorides	8.177184	11.288323	14.399461	5.741485e-13	TRUE
free.sulfur.dioxide	5.405034	7.655202	9.905369	1.297337e-11	TRUE
total.sulfur.dioxide	8.634538	12.041189	15.447840	2.138497e-12	TRUE
density	8.174147	11.335881	14.497615	1.054119e-12	TRUE
pH	6.517177	9.338854	12.160530	4.382196e-11	TRUE
sulphates	10.850121	14.984289	19.118457	6.064748e-13	TRUE
alcohol	16.558596	21.839066	27.119536	2.614475e-16	TRUE

```
# Variable Importance Plot
plot(vi)
```



```
# Confusion Matrix
# https://www.rdocumentation.org/packages/randomForestSRC/versions/3.1.0/topics/predict.rfsrc
randomForestSRC::predict.rfsrc(rf, rf_wine_test)
```

```

## Sample size of test (predict) data: 318
##           Number of grow trees: 500
## Average no. of grow terminal nodes: 252.772
##           Total no. of grow variables: 11
##           Resampling used to grow trees: swor
##           Resample size used to grow trees: 810
##           Analysis: RF-C
##           Family: class
##           Brier score: 0.07169764
##           Normalized Brier score: 0.516223
##           AUC: 0.84812834
##           Requested performance error: 0.32389937, 1, 1, 0.2406015, 0.26771654, 0.475, 1
##
## Confusion matrix:
##
##           predicted
## observed 3 4   5  6  7  8 class.error
##           3 0 0   4  1  0  0      1.0000
##           4 0 0   8  2  0  0      1.0000
##           5 0 0 101 32  0  0      0.2406
##           6 0 1  30 93  3  0      0.2677
##           7 0 0   0 19 21  0      0.4750
##           8 0 0   0  1  2  0      1.0000
##
##           Misclassification error: 0.3238994

```

Partial Least Squares

```

tctrl <- trainControl(method = "repeatedcv", repeats = 5, number = 10)

set.seed(4)
pls_wine <- train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data = wine_train,
  method = "pls",
  preProc = c("center", "scale", "BoxCox"),
  tunelength = 20,
  trControl = tctrl)

```

pls_wine

```

## Partial Least Squares
##
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5), Box-Cox transformation (5)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 1153, 1153, 1153, 1154, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      0.6406665  0.3615466  0.4955678

```

```

##      2      0.6401252  0.3626094  0.4954780
##      3      0.6401545  0.3625458  0.4949155
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 2.

# Add predicted values to new data frame
wine_test %>%
  mutate(predicted = predict(pls_wine, newdata = wine_test)) -> df3

# Summary of predicted interval
predict(pls_wine, newdata = wine_test, interval = "prediction") %>%
  summary()

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      4.399   5.305   5.632   5.655   6.012   6.791

# Confusion Matrix
# Convert predicted values to whole numbers, so they match target values
df3$predicted_int = as.integer(round(df3$predicted, digits = 0))

union3 <- union(df3$quality, df3$predicted_int)
table3 <- table(factor(df3$quality, union3), factor(df3$predicted_int, union3))

confusionMatrix(table3)

## Confusion Matrix and Statistics
##
##
##      5  6  4  7  8  3
##  5 84 48  1  0  0  0
##  6 35 87  0  5  0  0
##  4  7  3  0  0  0  0
##  7  0 31  0  9  0  0
##  8  0  3  0  0  0  0
##  3  4  0  1  0  0  0
##
## Overall Statistics
##
##              Accuracy : 0.566
##              95% CI   : (0.5096, 0.6212)
##      No Information Rate : 0.5409
##      P-Value [Acc > NIR] : 0.1995
##
##              Kappa   : 0.2854
##
##  Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: 5 Class: 6 Class: 4 Class: 7 Class: 8 Class: 3
## Sensitivity          0.6462   0.5058 0.000000   0.64286      NA      NA
## Specificity          0.7394   0.7260 0.968354   0.89803 0.990566  0.98428
## Pos Pred Value       0.6316   0.6850 0.000000   0.22500      NA      NA
## Neg Pred Value       0.7514   0.5550 0.993506   0.98201      NA      NA
## Prevalence           0.4088   0.5409 0.006289   0.04403 0.000000  0.00000


```

```
## Detection Rate      0.2642    0.2736 0.000000  0.02830 0.000000  0.00000
## Detection Prevalence 0.4182    0.3994 0.031447  0.12579 0.009434  0.01572
## Balanced Accuracy   0.6928    0.6159 0.484177  0.77044      NA      NA
```

```
# ROC plot
```

```
df3$predicted_int = round(as.numeric(as.character(df3$predicted))), digits = 0)
```

```
roc(df3$quality, df3$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)
```



```
##
```

```
## Call:
```

```
## roc.default(response = df3$quality, predictor = df3$predicted_int, percent = TRUE, plot = TRUE, ...)
```

```
##
```

```
## Data: df3$predicted_int in 5 controls (df3$quality 3) < 10 cases (df3$quality 4).
```

```
## Area under the curve: 72%
```

```
#
```

```
# modelName3 <- 'Partial Least Squares'
```

```
# roc3 <- multiclass.roc(df3$quality, df3$predicted_int)
```

```
# auc3 <- round(auc(df3$quality, df3$predicted_int), 4)
```

```
#
```


```
# ggroc(roc3, colours = 'red', size = 1) +
```

```
#   ggtitle(paste0(modelName3, ' - ROC Curve ', '(AUC = ', auc3, ')')) + theme_minimal()
```

```
# Scatter plot of predicted
```

```
ggplot(df3, aes(x = predicted, y = quality, colour = quality)) +
```

```
geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```



Mars Tuning

```
mars_wine <- earth(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data =wine_train)
```

```
mars_wine
```

```
## Selected 12 of 16 terms, and 5 of 5 predictors
## Termination condition: Reached nk 21
## Importance: alcohol, sulphates, volatile.acidity, total.sulfur.dioxide, ...
## Number of terms at each degree of interaction: 1 11 (additive model)
## GCV 0.4009379    RSS 495.3238    GRSq 0.3737305    RSq 0.3950735
```

```
summary(mars_wine)
```

```
## Call: earth(formula=quality~volatile.acidity+chlorides+total.sulfur.di...),
##           data=wine_train)
##
##
##               coefficients
## (Intercept)      47.707645
## h(0.44-volatile.acidity)  1.285938
## h(volatile.acidity-0.44) -0.806370
## h(chlorides-0.042)      20.942183
## h(chlorides-0.092)     -12.726397
## h(0.146-chlorides)      18.955464
## h(chlorides-0.226)     -10.995630
## h(total.sulfur.dioxide-9) -0.321577
## h(144-total.sulfur.dioxide) -0.318830
## h(total.sulfur.dioxide-144)  0.329662
```



```

## h(0.76-sulphates)          -2.012138
## h(12.5-alcohol)           -0.286173
##
## Selected 12 of 16 terms, and 5 of 5 predictors
## Termination condition: Reached nk 21
## Importance: alcohol, sulphates, volatile.acidity, total.sulfur.dioxide, ...
## Number of terms at each degree of interaction: 1 11 (additive model)
## GCV 0.4009379    RSS 495.3238    GRSq 0.3737305    RSq 0.3950735

preProc_Arguments = c("center", "scale")
marsGrid_wine = expand.grid(.degree=1:2, .nprune=2:38)

set.seed(4)

marsModel_wine = train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data =wine_train,
  method="earth",
  preProc=preProc_Arguments,
  tuneGrid=marsGrid_wine)

marsModel_wine

## Multivariate Adaptive Regression Spline
##
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1281, 1281, 1281, 1281, 1281, 1281, ...
## Resampling results across tuning parameters:
##
##  degree  nprune  RMSE      Rsquared  MAE
##  1         2    0.7014813  0.2419795  0.5549556
##  1         3    0.6684464  0.3108674  0.5216037
##  1         4    0.6485241  0.3512544  0.5051027
##  1         5    0.6512351  0.3469701  0.5048202
##  1         6    0.6504907  0.3485716  0.5022178
##  1         7    0.6507935  0.3481485  0.5022971
##  1         8    0.6488589  0.3515400  0.5017645
##  1         9    0.6503107  0.3492233  0.5024412
##  1        10    0.6533237  0.3449166  0.5036034
##  1        11    0.6554705  0.3422215  0.5043991
##  1        12    0.6552105  0.3429115  0.5041358
##  1        13    0.6555639  0.3424326  0.5046787
##  1        14    0.6558033  0.3422174  0.5046129
##  1        15    0.6563798  0.3412858  0.5050354
##  1        16    0.6563798  0.3412858  0.5050354
##  1        17    0.6563798  0.3412858  0.5050354
##  1        18    0.6563798  0.3412858  0.5050354
##  1        19    0.6563798  0.3412858  0.5050354
##  1        20    0.6563798  0.3412858  0.5050354
##  1        21    0.6563798  0.3412858  0.5050354
##  1        22    0.6563798  0.3412858  0.5050354
##  1        23    0.6563798  0.3412858  0.5050354

```

##	1	24	0.6563798	0.3412858	0.5050354
##	1	25	0.6563798	0.3412858	0.5050354
##	1	26	0.6563798	0.3412858	0.5050354
##	1	27	0.6563798	0.3412858	0.5050354
##	1	28	0.6563798	0.3412858	0.5050354
##	1	29	0.6563798	0.3412858	0.5050354
##	1	30	0.6563798	0.3412858	0.5050354
##	1	31	0.6563798	0.3412858	0.5050354
##	1	32	0.6563798	0.3412858	0.5050354
##	1	33	0.6563798	0.3412858	0.5050354
##	1	34	0.6563798	0.3412858	0.5050354
##	1	35	0.6563798	0.3412858	0.5050354
##	1	36	0.6563798	0.3412858	0.5050354
##	1	37	0.6563798	0.3412858	0.5050354
##	1	38	0.6563798	0.3412858	0.5050354
##	2	2	0.7014414	0.2419923	0.5548719
##	2	3	0.6702386	0.3075375	0.5230336
##	2	4	0.6524546	0.3440511	0.5064611
##	2	5	0.6462661	0.3563502	0.5004076
##	2	6	0.6452593	0.3592572	0.4978439
##	2	7	0.6465667	0.3567901	0.4989545
##	2	8	0.6494945	0.3521291	0.5005637
##	2	9	0.6505045	0.3507097	0.5006272
##	2	10	0.6530531	0.3462806	0.5010743
##	2	11	0.6560394	0.3415267	0.5022654
##	2	12	0.6560221	0.3419308	0.5026598
##	2	13	0.6570923	0.3406133	0.5033740
##	2	14	0.6586577	0.3389605	0.5040339
##	2	15	0.6595409	0.3376501	0.5044626
##	2	16	0.6614910	0.3343875	0.5057373
##	2	17	0.6618613	0.3339272	0.5059127
##	2	18	0.6618503	0.3340194	0.5060600
##	2	19	0.6620816	0.3337136	0.5061720
##	2	20	0.6621197	0.3336746	0.5061551
##	2	21	0.6621197	0.3336746	0.5061551
##	2	22	0.6621197	0.3336746	0.5061551
##	2	23	0.6621197	0.3336746	0.5061551
##	2	24	0.6621197	0.3336746	0.5061551
##	2	25	0.6621197	0.3336746	0.5061551
##	2	26	0.6621197	0.3336746	0.5061551
##	2	27	0.6621197	0.3336746	0.5061551
##	2	28	0.6621197	0.3336746	0.5061551
##	2	29	0.6621197	0.3336746	0.5061551
##	2	30	0.6621197	0.3336746	0.5061551
##	2	31	0.6621197	0.3336746	0.5061551
##	2	32	0.6621197	0.3336746	0.5061551
##	2	33	0.6621197	0.3336746	0.5061551
##	2	34	0.6621197	0.3336746	0.5061551
##	2	35	0.6621197	0.3336746	0.5061551
##	2	36	0.6621197	0.3336746	0.5061551
##	2	37	0.6621197	0.3336746	0.5061551
##	2	38	0.6621197	0.3336746	0.5061551

##

RMSE was used to select the optimal model using the smallest value.

```

## The final values used for the model were nprune = 6 and degree = 2.
# Add predicted values to new data frame
wine_test %>%
  mutate(predicted = predict(marsModel_wine, newdata = wine_test)) -> df4

# Summary of predicted interval
predict(marsModel_wine, newdata = wine_test, interval = "prediction") %>%
  summary()

##           y
## Min.      :3.974
## 1st Qu.:5.273
## Median :5.596
## Mean     :5.651
## 3rd Qu.:6.000
## Max.     :6.804

# Confusion Matrix
# Convert predicted values to whole numbers, so they match target values
df4$predicted_int = as.integer(round(df4$predicted, digits = 0))

union4 <- union(df4$quality, df4$predicted_int)
table4 <- table(factor(df4$quality, union4), factor(df4$predicted_int, union4))

confusionMatrix(table4)

## Confusion Matrix and Statistics
##
##
##      5  6  4  7  8  3
## 5 88 45  0  0  0  0
## 6 36 84  0  7  0  0
## 4  6  4  0  0  0  0
## 7  0 33  0  7  0  0
## 8  0  2  0  1  0  0
## 3  4  0  1  0  0  0
##
## Overall Statistics
##
##              Accuracy : 0.5629
##              95% CI   : (0.5064, 0.6182)
##    No Information Rate : 0.5283
##    P-Value [Acc > NIR] : 0.119
##
##              Kappa   : 0.2796
##
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: 5 Class: 6 Class: 4 Class: 7 Class: 8 Class: 3
## Sensitivity          0.6567   0.5000 0.000000  0.46667      NA      NA
## Specificity          0.7554   0.7133 0.968454  0.89109 0.990566  0.98428
## Pos Pred Value       0.6617   0.6614 0.000000  0.17500      NA      NA
## Neg Pred Value       0.7514   0.5602 0.996753  0.97122      NA      NA


```

```
## Prevalence          0.4214    0.5283 0.003145  0.04717 0.000000  0.00000
## Detection Rate      0.2767    0.2642 0.000000  0.02201 0.000000  0.00000
## Detection Prevalence 0.4182    0.3994 0.031447  0.12579 0.009434  0.01572
## Balanced Accuracy    0.7061    0.6067 0.484227  0.67888      NA      NA
```

```
# ROC plot
```

```
df4$predicted_int = round(as.numeric(as.character(df4$predicted))), digits = 0)
```

```
roc(df4$quality, df4$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)
```



```
##
```

```
## Call:
```

```
## roc.default(response = df4$quality, predictor = df4$predicted_int, percent = TRUE, plot = TRUE, ...)
```

```
##
```

```
## Data: df4$predicted_int in 5 controls (df4$quality 3) < 10 cases (df4$quality 4).
```

```
## Area under the curve: 76%
```

```
#
```

```
# modelName4 <- 'Mars Tuning'
```

```
# roc4 <- multiclass.roc(df4$quality, df4$predicted_int)
```

```
# auc4 <- round(auc(df4$quality, df4$predicted_int), 4)
```

```
#
```


```
# ggroc(roc4, colours = 'red', size = 1) +
```

```
# ggtitle(paste0(modelName4, ' - ROC Curve ', '(AUC = ', auc4, ', ')') + theme_minimal())
```

```
# Scatter plot of predicted
```

```
ggplot(df4, aes(x = predicted, y = quality, colour = quality)) +
```

```
geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```



KNN Neighbors

```
set.seed(4)

knn_wine <- train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data =wine_train,
  method = "knn",
  preProc = c("center", "scale"),
  tuneGrid = data.frame(.k = 1:50),
  trControl = trainControl(method = "cv"))

knn_wine
```

```
## k-Nearest Neighbors
##
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1153, 1153, 1153, 1154, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##  k  RMSE      Rsquared  MAE
##  1  0.7931263  0.2558058  0.4573932
##  2  0.7321340  0.2784365  0.5028836
##  3  0.7050894  0.2946256  0.5038454
##  4  0.6928700  0.3007317  0.5152042
```

```
##      5  0.6824992  0.3093344  0.5120538
##      6  0.6747969  0.3146795  0.5094303
##      7  0.6648928  0.3275608  0.5063056
##      8  0.6571997  0.3399876  0.5000693
##      9  0.6550702  0.3423538  0.5014153
##     10  0.6534921  0.3441642  0.5036200
##     11  0.6506971  0.3483598  0.5023332
##     12  0.6498040  0.3492250  0.5031206
##     13  0.6480800  0.3519414  0.5036606
##     14  0.6494254  0.3488478  0.5048960
##     15  0.6501741  0.3464884  0.5060485
##     16  0.6487798  0.3486218  0.5045870
##     17  0.6482825  0.3492570  0.5042702
##     18  0.6471091  0.3512579  0.5027609
##     19  0.6472611  0.3511440  0.5033137
##     20  0.6463476  0.3525807  0.5026752
##     21  0.6442849  0.3561775  0.5011751
##     22  0.6442341  0.3559528  0.5014678
##     23  0.6444608  0.3553065  0.5020588
##     24  0.6444540  0.3551886  0.5018186
##     25  0.6425414  0.3583346  0.5007683
##     26  0.6410799  0.3613750  0.4996561
##     27  0.6411065  0.3615645  0.5001382
##     28  0.6405786  0.3628237  0.4998135
##     29  0.6393510  0.3655226  0.4985531
##     30  0.6380821  0.3682214  0.4980229
##     31  0.6383581  0.3675952  0.4983101
##     32  0.6384969  0.3675095  0.4981278
##     33  0.6385071  0.3677068  0.4982906
##     34  0.6382805  0.3684733  0.4982302
##     35  0.6373637  0.3702965  0.4979635
##     36  0.6374273  0.3700186  0.4983200
##     37  0.6369241  0.3709937  0.4977353
##     38  0.6364424  0.3718126  0.4979826
##     39  0.6361429  0.3724760  0.4979734
##     40  0.6355148  0.3740180  0.4975481
##     41  0.6360932  0.3729431  0.4980664
##     42  0.6352480  0.3746541  0.4978058
##     43  0.6345464  0.3761934  0.4968382
##     44  0.6350254  0.3753575  0.4972071
##     45  0.6352044  0.3751202  0.4977586
##     46  0.6353060  0.3749478  0.4982237
##     47  0.6346586  0.3762757  0.4975144
##     48  0.6347369  0.3761214  0.4978120
##     49  0.6351326  0.3753620  0.4980585
##     50  0.6354252  0.3748959  0.4981426
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 43.
# Add predicted values to new data frame
wine_test %>%
  mutate(predicted = predict(knn_wine, newdata = wine_test)) -> df5
```

```

# Summary of predicted interval
predict(knn_wine, newdata = wine_test, interval = "prediction") %>%
  summary()

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      4.860  5.302   5.605   5.673   5.977   6.814

# Confusion Matrix
# Convert predicted values to whole numbers, so they match target values
df5$predicted_int = as.integer(round(df5$predicted, digits = 0))

union5 <- union(df5$quality, df5$predicted_int)
table5 <- table(factor(df5$quality, union5), factor(df5$predicted_int, union5))


confusionMatrix(table5)

## Confusion Matrix and Statistics
##
##
##      5  6  4  7  8  3
##  5 87 46  0  0  0  0
##  6 32 88  0  7  0  0
##  4  6  4  0  0  0  0
##  7  2 30  0  8  0  0
##  8  0  1  0  2  0  0
##  3  5  0  0  0  0  0
##
## Overall Statistics
##
##              Accuracy : 0.5755
##              95% CI : (0.5191, 0.6304)
##      No Information Rate : 0.5314
##      P-Value [Acc > NIR] : 0.06436
##
##              Kappa : 0.3011
##
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: 5 Class: 6 Class: 4 Class: 7 Class: 8 Class: 3
## Sensitivity          0.6591  0.5207      NA  0.47059      NA      NA
## Specificity          0.7527  0.7383  0.96855  0.89369 0.990566  0.98428
## Pos Pred Value       0.6541  0.6929      NA  0.20000      NA      NA
## Neg Pred Value       0.7568  0.5759      NA  0.96763      NA      NA
## Prevalence           0.4151  0.5314  0.00000  0.05346 0.000000  0.00000
## Detection Rate       0.2736  0.2767  0.00000  0.02516 0.000000  0.00000
## Detection Prevalence 0.4182  0.3994  0.03145  0.12579 0.009434  0.01572
## Balanced Accuracy     0.7059  0.6295      NA  0.68214      NA      NA

# ROC plot
df5$predicted_int = round(as.numeric(as.character(df5$predicted)), digits = 0)

roc(df5$quality, df5$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)


```



```
##
## Call:
## roc.default(response = df5$quality, predictor = df5$predicted_int, percent = TRUE, plot = TRUE, ...)
##
## Data: df5$predicted_int in 5 controls (df5$quality 3) < 10 cases (df5$quality 4).
## Area under the curve: 70%
```

```
#
# modelName5 <- 'KNN'
# roc5 <- multiclass.roc(df5$quality, df5$predicted_int)
# auc5 <- round(auc(df5$quality, df5$predicted_int), 4)
#
# ggroc(roc5, colours = 'red', size = 1) +
#   ggtitle(paste0(modelName5, ' - ROC Curve ', '(AUC = ', auc5, ')')) + theme_minimal()

# Scatter plot of predicted
ggplot(df5, aes(x = predicted, y = quality, colour = quality)) +
  geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```

SVM

```
set.seed(4)

svmTune <- train(quality ~ volatile.acidity + sulphates + alcohol, data = rf_wine_train, # using the su
                method = "svmRadial",
                preProc = c("center", "scale"),
                tuneLength= 5,
                trControl = trainControl(method = "cv"))

svmTune

## Support Vector Machines with Radial Basis Function Kernel
##
## 1281 samples
##    3 predictor
##    6 classes: '3', '4', '5', '6', '7', '8'
##
## Pre-processing: centered (3), scaled (3)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1153, 1154, 1153, 1152, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##    C      Accuracy  Kappa
## 0.25 0.5910090 0.3127878
## 0.50 0.5925714 0.3185638
## 1.00 0.5933528 0.3215787
## 2.00 0.5870903 0.3128096
```

```

##    4.00  0.5901849  0.3205236
##
## Tuning parameter 'sigma' was held constant at a value of 0.6015525
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were sigma = 0.6015525 and C = 1.

# Add predicted values to new data frame
wine_test %>%
  mutate(predicted = predict(svmTune, newdata = wine_test)) -> df6

# Summary of predicted interval
predict(svmTune, newdata = wine_test, interval = "prediction") %>%
  summary()

##    3    4    5    6    7    8
##    0    0 152 152  14    0


# Confusion Matrix
confusionMatrix(table(df6$quality, df6$predicted))

## Confusion Matrix and Statistics
##
##
##      3  4  5  6  7  8
##  3  0  0  5  0  0  0
##  4  0  0  7  3  0  0
##  5  0  0 96 37  0  0
##  6  0  0 43 80  4  0
##  7  0  0  1 30  9  0
##  8  0  0  0  2  1  0
##
## Overall Statistics
##
##               Accuracy : 0.5818
##               95% CI : (0.5254, 0.6366)
##      No Information Rate : 0.478
##      P-Value [Acc > NIR] : 0.000131
##
##               Kappa : 0.3072
##
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: 3 Class: 4 Class: 5 Class: 6 Class: 7 Class: 8
## Sensitivity           NA      NA  0.6316  0.5263  0.64286      NA
## Specificity          0.98428 0.96855  0.7771  0.7169  0.89803 0.990566
## Pos Pred Value        NA      NA  0.7218  0.6299  0.22500      NA
## Neg Pred Value        NA      NA  0.6973  0.6230  0.98201      NA
## Prevalence            0.00000 0.00000  0.4780  0.4780  0.04403 0.000000
## Detection Rate        0.00000 0.00000  0.3019  0.2516  0.02830 0.000000
## Detection Prevalence  0.01572 0.03145  0.4182  0.3994  0.12579 0.009434
## Balanced Accuracy      NA      NA  0.7043  0.6216  0.77044      NA

# ROC plot
df6$predicted_int = round(as.numeric(as.character(df6$predicted)), digits = 0)

```


```
roc(df6$quality, df6$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)
```



```
##
## Call:
## roc.default(response = df6$quality, predictor = df6$predicted_int, percent = TRUE, plot = TRUE,
##
## Data: df6$predicted_int in 5 controls (df6$quality 3) < 10 cases (df6$quality 4).
## Area under the curve: 65%

#
# modelName6 <- 'SVM'
# roc6 <- multiclass.roc(df6$quality, df6$predicted_int)
# auc6 <- round(auc(df6$quality, df6$predicted_int), 4)
#
# ggroc(roc6, colours = 'red', size = 1) +
#   ggtitle(paste0(modelName6, ' - ROC Curve ', '(AUC = ', auc6, ')')) + theme_minimal()

# Scatter plot of predicted
ggplot(df6, aes(x = predicted, y = quality, colour = quality)) +
  geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```



Penalized Logistic Regression Tuning

```
#tuning parameters, alpha is associated with the ridge(0) versus lasso regression(1)
glmnetGrid <- expand.grid(alpha = c(0, .1, .2, .4, .6, .8, 1),
                          lambda = seq(.01, .2, length = 5))
glmnetTune <- train(quality ~ ., data = rf_wine_train, # using the subset data as used in random forest,
                    method = "glmnet",
                    tuneGrid = glmnetGrid,
                    preProc = c("center", "scale"),
                    trControl = trainControl(method = "cv"))

glmnetTune
```

```
## glmnet
##
## 1281 samples
## 11 predictor
## 6 classes: '3', '4', '5', '6', '7', '8'
##
## Pre-processing: centered (11), scaled (11)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1154, 1154, 1152, 1152, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##  alpha  lambda  Accuracy  Kappa
##  0.0    0.0100  0.6010676  0.33503660
##  0.0    0.0575  0.5878469  0.30509871
##  0.0    0.1050  0.5864007  0.29798813
##  0.0    0.1525  0.5856007  0.29432135
```

```
## 0.0 0.2000 0.5832932 0.28882488
## 0.1 0.0100 0.6034296 0.34232307
## 0.1 0.0575 0.5918024 0.30996661
## 0.1 0.1050 0.5832260 0.28973398
## 0.1 0.1525 0.5809246 0.28441948
## 0.1 0.2000 0.5777625 0.27857370
## 0.2 0.0100 0.6018732 0.33995345
## 0.2 0.0575 0.5839765 0.29403575
## 0.2 0.1050 0.5824382 0.28718508
## 0.2 0.1525 0.5739230 0.27193416
## 0.2 0.2000 0.5739167 0.27137366
## 0.4 0.0100 0.6049557 0.34275408
## 0.4 0.0575 0.5778233 0.28107401
## 0.4 0.1050 0.5731295 0.27041928
## 0.4 0.1525 0.5676840 0.26012440
## 0.4 0.2000 0.5645157 0.25395588
## 0.6 0.0100 0.6002494 0.33458111
## 0.6 0.0575 0.5723418 0.27043675
## 0.6 0.1050 0.5661034 0.25782801
## 0.6 0.1525 0.5613780 0.24867794
## 0.6 0.2000 0.5566783 0.23925609
## 0.8 0.0100 0.6010799 0.33417339
## 0.8 0.0575 0.5723667 0.26954382
## 0.8 0.1050 0.5637524 0.25326580
## 0.8 0.1525 0.5558909 0.23814660
## 0.8 0.2000 0.5472657 0.22180362
## 1.0 0.0100 0.5987421 0.32925132
## 1.0 0.0575 0.5739351 0.27169914
## 1.0 0.1050 0.5605787 0.24682498
## 1.0 0.1525 0.5559031 0.23773006
## 1.0 0.2000 0.4646061 0.06934246
##
```

```
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were alpha = 0.4 and lambda = 0.01.
```

```
# Add predicted values to new data frame
```

```
wine_test %>%
```

```
  mutate(predicted = predict(glmnTune, newdata = wine_test)) -> df7
```

```
# Summary of predicted interval
```

```
predict(glmnTune, newdata = wine_test, interval = "prediction") %>%
  summary()
```

```
## 3 4 5 6 7 8
## 0 1 146 156 15 0
```

```
# Confusion Matrix
```

```
confusionMatrix(table(df7$quality, df7$predicted))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##
```

```
##      3 4 5 6 7 8
```

```
## 3 0 1 4 0 0 0
```

```
## 4 0 0 7 3 0 0
```

```
## 5 0 0 91 42 0 0
## 6 0 0 43 76 8 0
## 7 0 0 1 33 6 0
## 8 0 0 0 2 1 0
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
## Accuracy : 0.544
```

```
## 95% CI : (0.4875, 0.5997)
```

```
## No Information Rate : 0.4906
```

```
## P-Value [Acc > NIR] : 0.03207
```

```
##
```

```
## Kappa : 0.2476
```

```
##
```

```
## McNemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
```

```
## Class: 3 Class: 4 Class: 5 Class: 6 Class: 7 Class: 8
```

```
## Sensitivity NA 0.000000 0.6233 0.4872 0.40000 NA
```

```
## Specificity 0.98428 0.968454 0.7558 0.6852 0.88779 0.990566
```

```
## Pos Pred Value NA 0.000000 0.6842 0.5984 0.15000 NA
```

```
## Neg Pred Value NA 0.996753 0.7027 0.5812 0.96763 NA
```

```
## Prevalence 0.00000 0.003145 0.4591 0.4906 0.04717 0.000000
```

```
## Detection Rate 0.00000 0.000000 0.2862 0.2390 0.01887 0.000000
```


```
## Detection Prevalence 0.01572 0.031447 0.4182 0.3994 0.12579 0.009434
```

```
## Balanced Accuracy NA 0.484227 0.6896 0.5862 0.64389 NA
```

```
# ROC plot
```

```
df7$predicted_int = round(as.numeric(as.character(df7$predicted))), digits = 0)
```

```
roc(df7$quality, df7$predicted_int, plot = TRUE, legacy.axes = TRUE, percent = TRUE, print.auc = TRUE)
```



```
##
## Call:
## roc.default(response = df7$quality, predictor = df7$predicted_int, percent = TRUE, plot = TRUE, ...)
##
## Data: df7$predicted_int in 5 controls (df7$quality 3) < 10 cases (df7$quality 4).
## Area under the curve: 72%
```

```
#
# modelName7 <- 'Penalized Logistic Regression Tuning'
# roc7 <- roc(df7$quality, df7$predicted_int)
# auc7 <- round(auc(roc7), 4)
#
# ggroc(roc7, colours = 'red', size = 1) +
#   ggtitle(paste0(modelName7, ' - ROC Curve ', '(AUC = ', auc7, ')')) + theme_minimal()

# Scatter plot of predicted
ggplot(df7, aes(x = predicted, y = quality, colour = quality)) +
  geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```

