

ADS 503 - Team 7

Summer Purschke, Jacqueline Urenda, Oscar Gil

06/12/2022

```
# R Libraries
library(caret)
library(AppliedPredictiveModeling)
library(Hmisc)
library(dplyr)
library(tidyverse)
library(ggplot2)
library(corrplot)
library(MASS)
library(ISLR)
library(rpart)
library(partykit)
library(randomForestSRC)
library(earth)
library(MARSS)
library(e1071)
library(summarytools)
```

Load the Red Wine Quality data set from GitHub - data set copied from Kaggle and imported into GitHub.

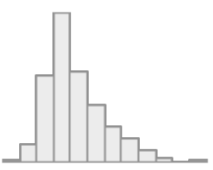
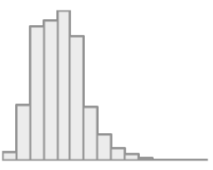
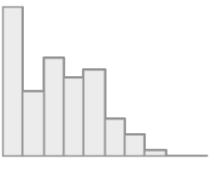
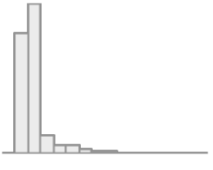
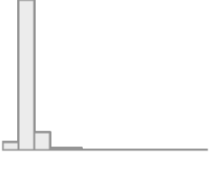
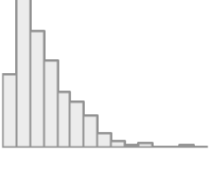
```
wine <- read.csv(
  url("https://raw.githubusercontent.com/OscarG-DataSci/ADS503/main/winequality-red.csv"),
  , header = TRUE)
```

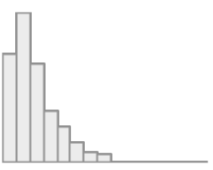
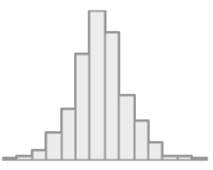
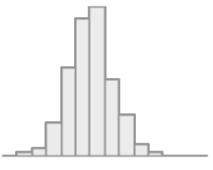
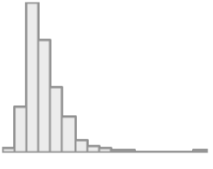
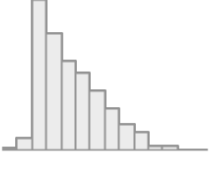
Data Summary


```
dfSummary(wine,
  plain.ascii = FALSE,
  style       = "grid",
  graph.magnif = 0.75,
  valid.col   = FALSE,
  tmp.img.dir  = "/tmp")
```

Data Frame Summary

wine Dimensions: 1599 x 12
Duplicates: 240

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
1	fixed.acidity [numeric]	Mean (sd) : 8.3 (1.7) min < med < max: 4.6 < 7.9 < 15.9 IQR (CV) : 2.1 (0.2)	96 distinct values		0 (0.0%)
2	volatile.acidity [numeric]	Mean (sd) : 0.5 (0.2) min < med < max: 0.1 < 0.5 < 1.6 IQR (CV) : 0.2 (0.3)	143 distinct values		0 (0.0%)
3	citric.acid [numeric]	Mean (sd) : 0.3 (0.2) min < med < max: 0 < 0.3 < 1 IQR (CV) : 0.3 (0.7)	80 distinct values		0 (0.0%)
4	residual.sugar [numeric]	Mean (sd) : 2.5 (1.4) min < med < max: 0.9 < 2.2 < 15.5 IQR (CV) : 0.7 (0.6)	91 distinct values		0 (0.0%)
5	chlorides [numeric]	Mean (sd) : 0.1 (0) min < med < max: 0 < 0.1 < 0.6 IQR (CV) : 0 (0.5)	153 distinct values		0 (0.0%)
6	free.sulfur.dioxide [numeric]	Mean (sd) : 15.9 (10.5) min < med < max: 1 < 14 < 72 IQR (CV) : 14 (0.7)	60 distinct values		0 (0.0%)

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
7	total.sulfur.dioxide [numeric]	Mean (sd) : 46.5 (32.9) min < med < max: 6 < 38 < 289 IQR (CV) : 40 (0.7)	144 distinct values		0 (0.0%)
8	density [numeric]	Mean (sd) : 1 (0) min < med < max: 1 < 1 < 1 IQR (CV) : 0 (0)	436 distinct values		0 (0.0%)
9	pH [numeric]	Mean (sd) : 3.3 (0.2) min < med < max: 2.7 < 3.3 < 4 IQR (CV) : 0.2 (0)	89 distinct values		0 (0.0%)
10	sulphates [numeric]	Mean (sd) : 0.7 (0.2) min < med < max: 0.3 < 0.6 < 2 IQR (CV) : 0.2 (0.3)	96 distinct values		0 (0.0%)
11	alcohol [numeric]	Mean (sd) : 10.4 (1.1) min < med < max: 8.4 < 10.2 < 14.9 IQR (CV) : 1.6 (0.1)	65 distinct values		0 (0.0%)

No	Variable	Stats / Values	Freqs (% of Valid)	Graph	Missing
					
12	quality [integer]	Mean (sd) : 5.6 (0.8) min < med < max: 3 < 6 < 8 IQR (CV) : 1 (0.1)	3 : 10 (0.6%) 4 : 53 (3.3%) 5 : 681 (42.6%) 6 : 638 (39.9%) 7 : 199 (12.4%) 8 : 18 (1.1%)		0 (0.0%)

Pre-processing

```
par(mar=c(1,1,1,1)) # to fix boxplot knit processing issues
```

```
# Create new variable, for quality values, split by half (0, 1)
wine$quality_target <- ifelse( wine$quality <= 5, 0, 1)
```

```
# Mean of new variable is at 0.5347 (close enough to 50% to maintain balance)
summary(wine$quality_target)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000  0.0000  1.0000  0.5347  1.0000  1.0000
```

```
# Check for missing values in data set
wine %>% na.omit() %>% count() # there are no missing values
```

```
##      n
## 1 1599
```

```
# Removing outliers for residual sugar:
```

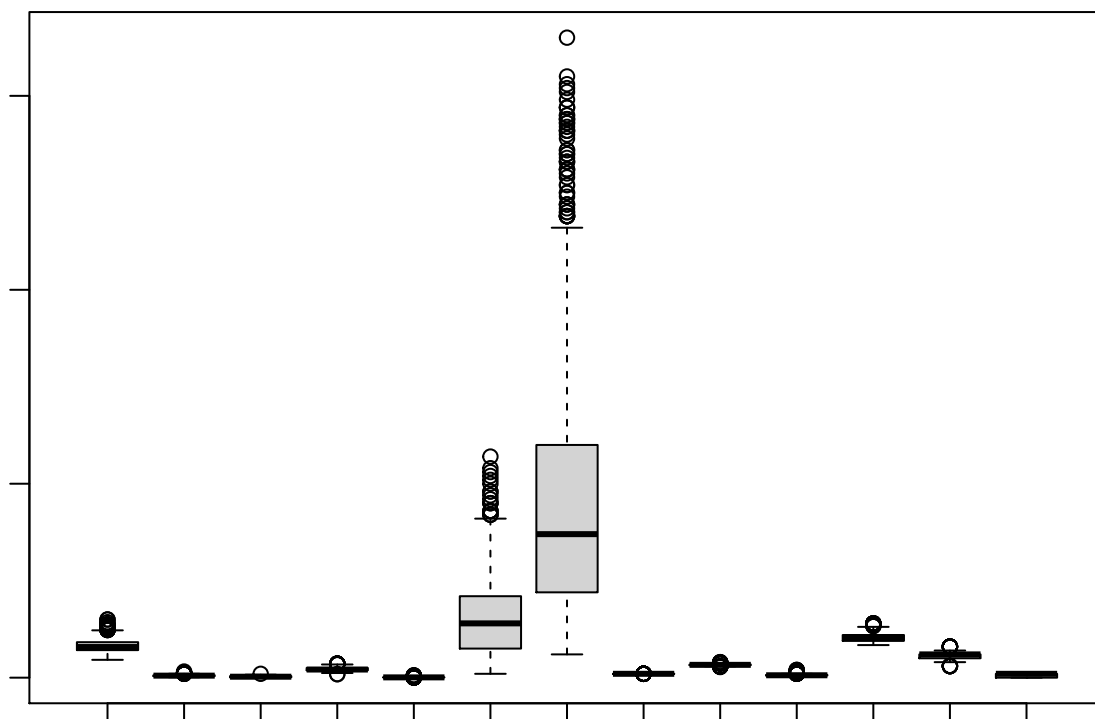
```
Q <- quantile(wine$residual.sugar, probs=c(.25, .75), na.rm = FALSE)
```

```
iqr_rs <- IQR(wine$residual.sugar)
```

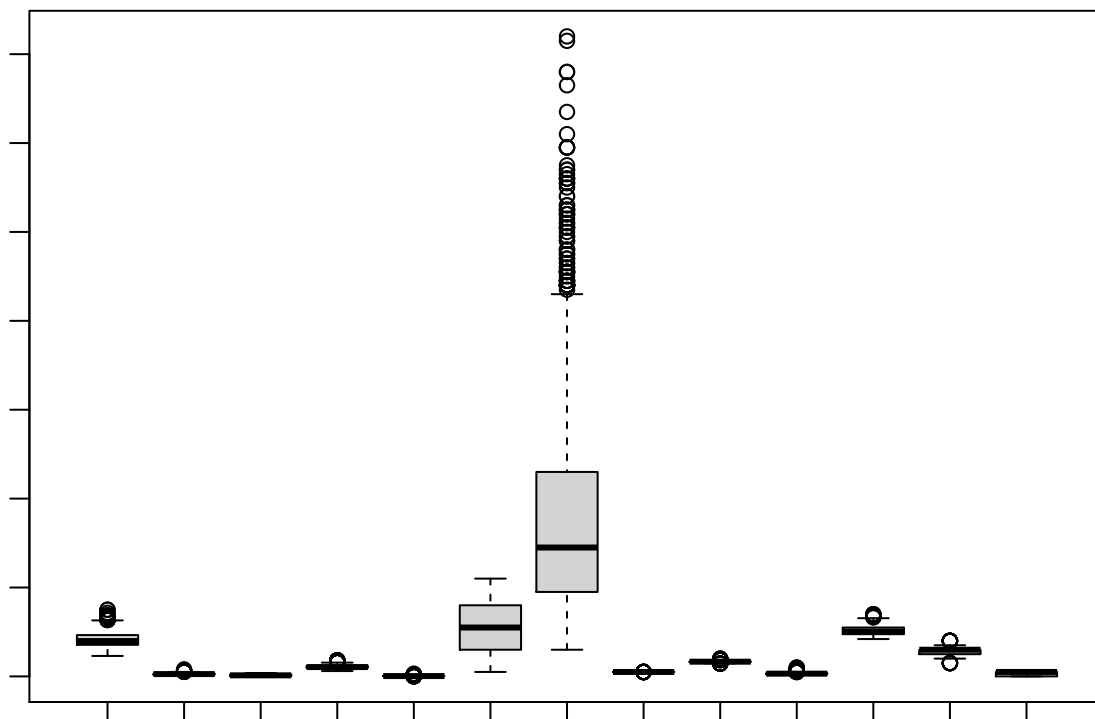
```
up_rs <- Q[2]+1.5*iqr_rs # Upper Range
```

```
low_rs <- Q[1]-1.5*iqr_rs # Lower Range
```

```
eliminated_rs <- subset(wine, wine$residual.sugar > (Q[1] - 1.5*iqr_rs) & wine$residual.sugar < (Q[2]+1.5*iqr_rs))
boxplot(eliminated_rs)
```



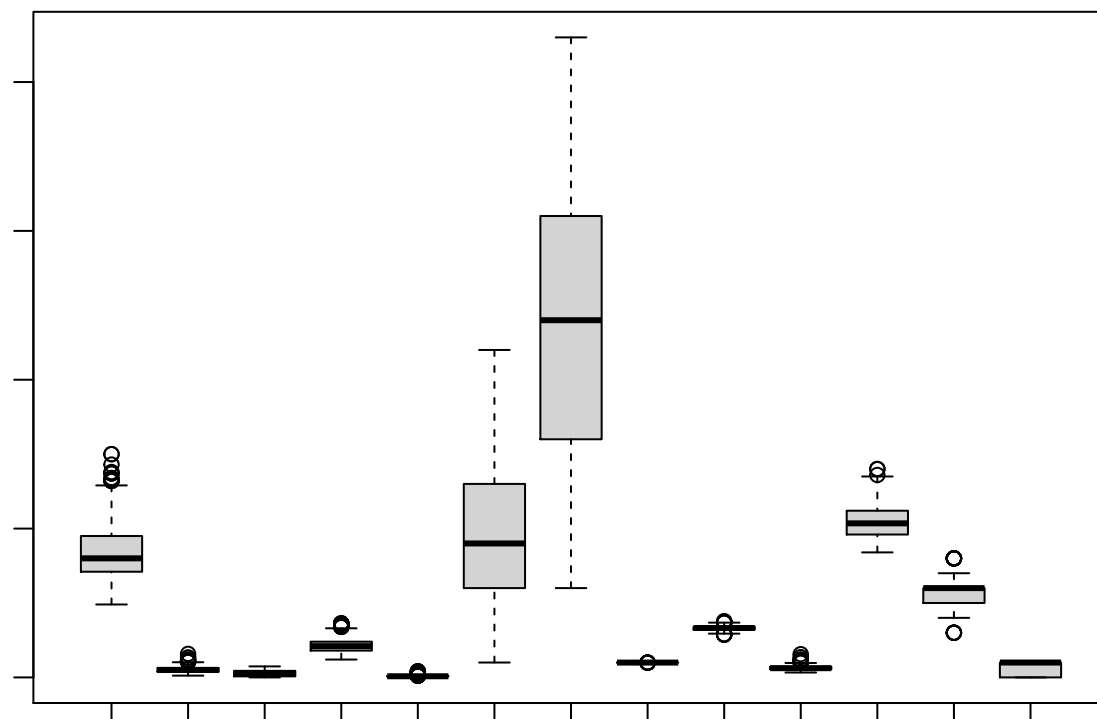
```
#Removing outliers for free.sulfur.dioxide:
Q2 <- quantile(wine$free.sulfur.dioxide, probs=c(.25, .75), na.rm = FALSE)
iqr_fs <- IQR(eliminated_rs$free.sulfur.dioxide)
up_fs <- Q2[2]+1.5*iqr_fs # Upper Range
low_fs <- Q2[1]-1.5*iqr_fs # Lower Range
eliminated_fs <- subset(eliminated_rs, eliminated_rs$free.sulfur.dioxide > (Q[1] - 1.5*iqr_fs) & eliminated_rs$free.sulfur.dioxide < (Q[2] + 1.5*iqr_fs))
boxplot(eliminated_fs)
```



```

#Removing outliers for total.sulfur.dioxide:
Q3 <- quantile(wine$total.sulfur.dioxide, probs=c(.25, .75), na.rm = FALSE)
iqr_ts <- IQR(eliminated_fs$total.sulfur.dioxide)
up_ts <- Q3[2]+1.5*iqr_ts # Upper Range
low_ts <- Q3[1]-1.5*iqr_ts # Lower Range
eliminated_ts <- subset(eliminated_fs, eliminated_fs$total.sulfur.dioxide > (Q[1] - 1.5*iqr_ts) & eliminated_ts < low_ts)
boxplot(eliminated_ts)

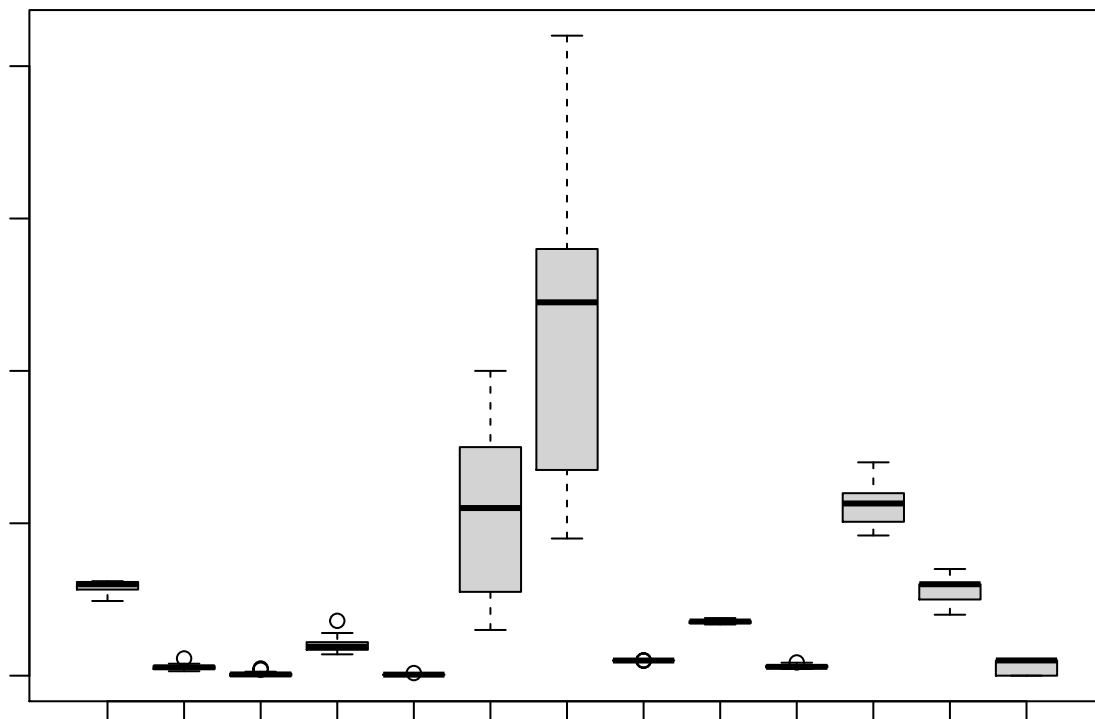
```



```

#Removing outliers for fixed.acidity:
Q4 <- quantile(wine$fixed.acidity, probs=c(.25, .75), na.rm = FALSE)
iqr_fa <- IQR(eliminated_ts$fixed.acidity)
up_fa <- Q[2]+1.5*iqr_fa # Upper Range
low_fa <- Q[1]-1.5*iqr_fa # Lower Range
eliminated_fa <- subset(eliminated_ts, eliminated_ts$fixed.acidity > (Q[1] - 1.5*iqr_fa) & eliminated_ts < low_fa)
boxplot(eliminated_fa)

```



```
new_wine_data <- eliminated_fa
```

```
# Removing outliers reduced dimension of data set from 1599 observations to 48
```

```
# team opted not to use new_wine_data and keep outlier data
```

```
dim(new_wine_data)
```

```
## [1] 48 13
```

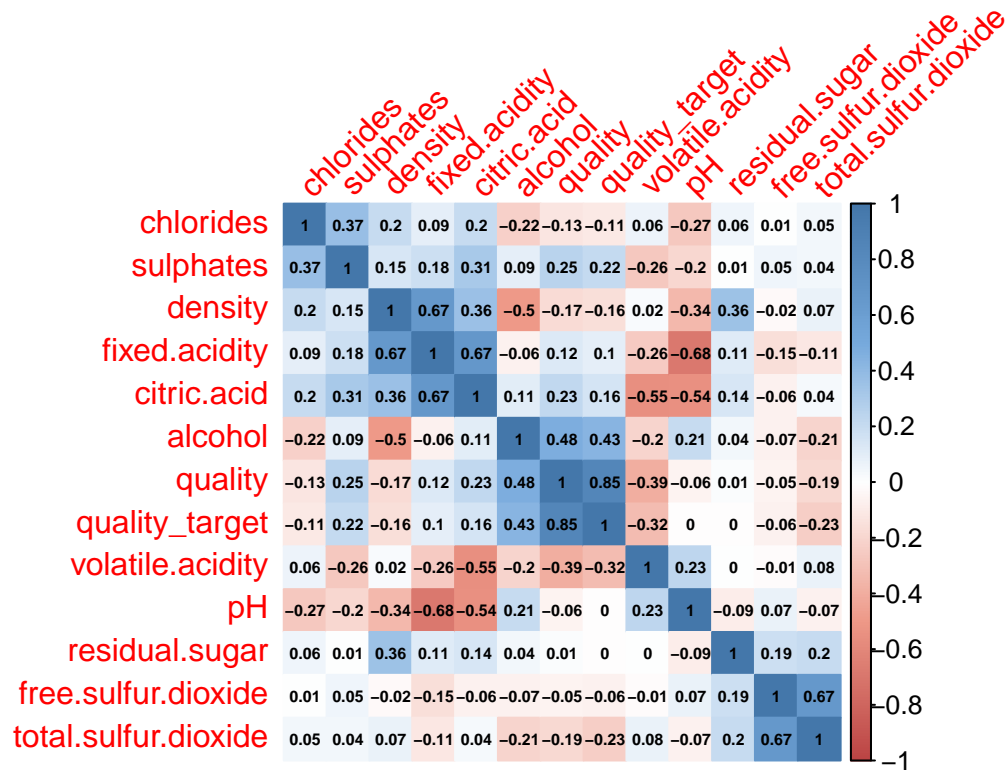
```
# Correlation Matrix
```

```
cor <- cor(wine)
```

```
# Colors for Correlation Matrix
```

```
colors <- colorRampPalette(c("#BB4444", "#EE9988", "#FFFFFF", "#77AADD", "#4477AA"))
```

```
corrplot(cor, order="hclust", method = "color", addCoef.col = "black",  
          , tl.srt = 45, number.cex = 0.47, col=colors(200))
```




```
# Cutoff Correlation features
cutoffCorr <- findCorrelation(cor, cutoff = .8)
cutoffCorrFeatures <- wine[, -cutoffCorr]

# Train and Test split
wine_split <- createDataPartition(wine$quality, p = .8, list = FALSE)
wine_train <- wine[ wine_split,]
wine_test  <- wine[-wine_split,]

# Transform Train Data
train_trans <- preProcess(wine_train, method = c("center", "scale"))
train_transformed <- predict(train_trans, wine_train)

# Transform Test Data
test_trans <- preProcess(wine_test, method = c("center", "scale"))
test_transformed <- predict(test_trans, wine_test)

# Boxplot of transformed train data
boxplot(train_transformed, horizontal = TRUE, las = 2, cex.axis = .65, cex.lab = 7)
```

Logistic Regression Model

Cutoff Correlation string to copy + paste into feature area of model

```
subset(cutoffCorrFeatures, select = -c(quality_target)) %>%
  colnames() %>%
  paste0(collapse = " + ")
```

```
## [1] "fixed.acidity + volatile.acidity + citric.acid + residual.sugar + chlorides + free.sulfur.dioxide"
```

```
set.seed(4)
```

Model using "quality_target" as target variable

```
lmodel1 <- lm(quality_target ~ volatile.acidity + sulphates + alcohol, data = train_transformed)
```

```
summary(lmodel1)
```

```
##
```

```
## Call:
```

```
## lm(formula = quality_target ~ volatile.acidity + sulphates +
```

```
##     alcohol, data = train_transformed)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -2.09174 -0.69098 -0.05791  0.76874  1.95454
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
```

```
## (Intercept)   -4.401e-15  2.373e-02   0.000      1
```

```
## volatile.acidity -2.210e-01  2.504e-02 -8.824 < 2e-16 ***
```

```
## sulphates      1.349e-01  2.466e-02  5.469 5.44e-08 ***
```

```
## alcohol          3.890e-01  2.428e-02  16.018  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8492 on 1277 degrees of freedom
## Multiple R-squared:  0.2806, Adjusted R-squared:  0.2789
## F-statistic: 166 on 3 and 1277 DF, p-value: < 2.2e-16

# Model using "quality" as target variable
lmodel2 <- lm(quality~ volatile.acidity + sulphates + alcohol, data = train_transformed)

summary(lmodel2)


##
## Call:
## lm(formula = quality ~ volatile.acidity + sulphates + alcohol,
##     data = train_transformed)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.84361 -0.47510 -0.07854  0.55916  2.71337
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5.081e-16  2.260e-02   0.000      1
## volatile.acidity -2.829e-01  2.385e-02 -11.861 < 2e-16 ***
## sulphates       1.477e-01  2.349e-02   6.286 4.45e-10 ***
## alcohol         4.043e-01  2.313e-02  17.478 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.8089 on 1277 degrees of freedom
## Multiple R-squared:  0.3471, Adjusted R-squared:  0.3456
## F-statistic: 226.3 on 3 and 1277 DF, p-value: < 2.2e-16

# Add predicted values to new data frame
wine_test %>%
  mutate(predicted = predict(lmodel2, newdata = test_transformed)) -> df

# Summary of predicted interval
predict(lmodel2, newdata = test_transformed, interval = "prediction") %>%
  summary()

##           fit           lwr           upr
## Min.      :-1.41473   Min.      :-3.0079   Min.      :0.1785
## 1st Qu.: -0.43915   1st Qu.: -2.0279   1st Qu.: 1.1495
## Median : -0.07877   Median : -1.6671   Median : 1.5095
## Mean     : 0.00000   Mean     : -1.5895   Mean     : 1.5895
## 3rd Qu.: 0.36193   3rd Qu.: -1.2262   3rd Qu.: 1.9501
## Max.     : 2.14449   Max.     : 0.5457   Max.     : 3.7433

# Scatter plot of predicted
ggplot(df, aes(x = predicted, y = quality, colour = quality)) +
  geom_point(alpha = 0.3, position = position_jitter()) + stat_smooth()
```



The scatter plot supports the summary of the predicted interval, in the ranges of the fit, lower, and upper ranges. The R-squared value of 0.3283 of the model, indicates that this information can be predicted 33% of the time, with the data available, for the variance of the information.

CART

```
set.seed(4)
# Subset both train and test sets, to exclude "quality_target"
subset(train_transformed, select = -c(quality_target)) -> rf_wine_train
subset(test_transformed, select = -c(quality_target)) -> rf_wine_test



rPartTree <- rpart(quality ~ ., data = rf_wine_train)

rpartTree2 <- as.party(rPartTree)

# R-Squared plot
par(mfrow=c(1,2))
rsq.rpart(rPartTree)

##
## Regression tree:
## rpart(formula = quality ~ ., data = rf_wine_train)
##
## Variables actually used in tree construction:
```

```
## [1] alcohol          sulphates          total.sulfur.dioxide
## [4] volatile.acidity
##
## Root node error: 1280/1281 = 0.99922
##
## n= 1281
##
##      CP nsplit rel error  xerror  xstd
## 1  0.181052    0  1.00000 1.00110 0.041762
## 2  0.051414    1  0.81895 0.82072 0.040061
## 3  0.031271    2  0.76753 0.79048 0.037492
## 4  0.026076    3  0.73626 0.78981 0.036679
## 5  0.024991    4  0.71019 0.78510 0.036738
## 6  0.020343    5  0.68520 0.76959 0.035964
## 7  0.015736    6  0.66485 0.73676 0.034727
## 8  0.012461    7  0.64912 0.72182 0.033440
## 9  0.010909    8  0.63666 0.72145 0.033255
## 10 0.010000   10  0.61484 0.71444 0.032724
```

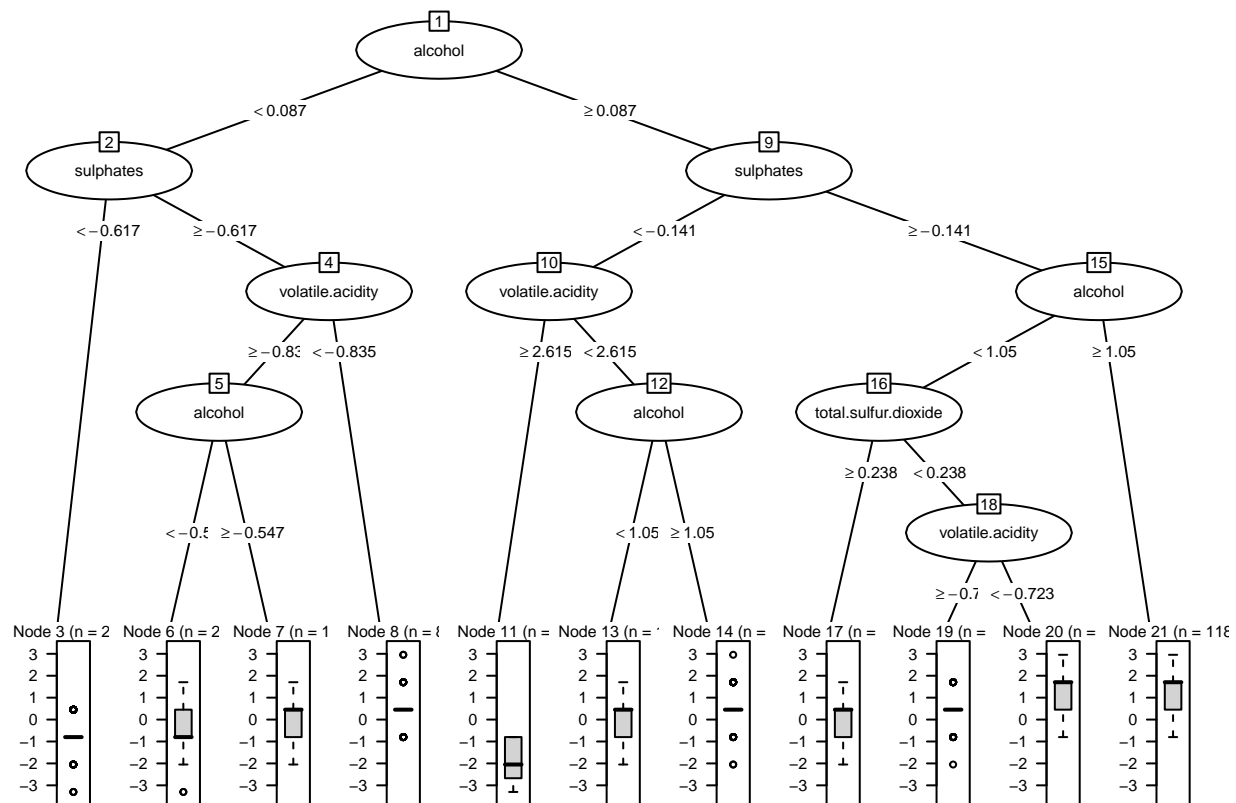


```
# Results
rpartTree2
```

```
##
## Model formula:
## quality ~ fixed.acidity + volatile.acidity + citric.acid + residual.sugar +
## chlorides + free.sulfur.dioxide + total.sulfur.dioxide +
## density + pH + sulphates + alcohol
##
## Fitted party:
## [1] root
## | [2] alcohol < 0.08723
## | | [3] sulphates < -0.61681: -0.666 (n = 254, err = 106.7)
## | | [4] sulphates >= -0.61681
```

```
## |   |   |   [5] volatile.acidity >= -0.83496
## |   |   |   [6] alcohol < -0.54685: -0.438 (n = 274, err = 138.1)
## |   |   |   [7] alcohol >= -0.54685: -0.050 (n = 172, err = 124.2)
## |   |   |   [8] volatile.acidity < -0.83496: 0.391 (n = 82, err = 56.1)
## |   [9] alcohol >= 0.08723
## |   |   [10] sulphates < -0.14052
## |   |   |   [11] volatile.acidity >= 2.61549: -1.894 (n = 8, err = 7.6)
## |   |   |   [12] volatile.acidity < 2.61549
## |   |   |   [13] alcohol < 1.05011: -0.082 (n = 117, err = 101.2)
## |   |   |   [14] alcohol >= 1.05011: 0.553 (n = 87, err = 66.4)
## |   |   [15] sulphates >= -0.14052
## |   |   |   [16] alcohol < 1.05011
## |   |   |   [17] total.sulfur.dioxide >= 0.23772: 0.099 (n = 39, err = 31.1)
## |   |   |   [18] total.sulfur.dioxide < 0.23772
## |   |   |   |   [19] volatile.acidity >= -0.72275: 0.435 (n = 73, err = 42.2)
## |   |   |   |   [20] volatile.acidity < -0.72275: 1.133 (n = 57, err = 44.1)
## |   |   [21] alcohol >= 1.05011: 1.205 (n = 118, err = 69.3)
##
## Number of inner nodes:    10
## Number of terminal nodes: 11
```

```
plot(rpartTree2, gp = gpar(fontsize=6))
```



Random Forest

```
set.seed(4)
```

```
rf <- rfsrc(quality ~ ., data = rf_wine_train)

print(rf)


##                      Sample size: 1281
##                      Number of trees: 500
##          Forest terminal node size: 5
##      Average no. of terminal nodes: 151.684
## No. of variables tried at each split: 4
##          Total no. of variables: 11
##      Resampling used to grow trees: swor
##      Resample size used to grow trees: 810
##                      Analysis: RF-R
##                      Family: regr
##                      Splitting rule: mse
##                      (OOB) R squared: 0.46705408
##      (OOB) Requested performance error: 0.53294592
```

```
# Variable Importance
vi <- subsample(rf, verbose = FALSE)

extract.subsample(vi)$var.jk.sel.Z
```

	lower	mean	upper	pvalue	signif
## fixed.acidity	28.95938	42.59931	56.23924	4.642868e-10	TRUE
## volatile.acidity	77.13856	99.97660	122.81465	4.743454e-18	TRUE
## citric.acid	29.83193	42.15590	54.47988	1.011647e-11	TRUE
## residual.sugar	45.31073	62.90817	80.50560	1.221103e-12	TRUE
## chlorides	41.57913	57.71305	73.84698	1.182815e-12	TRUE
## free.sulfur.dioxide	17.31180	25.16301	33.01422	1.674946e-10	TRUE
## total.sulfur.dioxide	31.37751	41.08702	50.79654	5.482936e-17	TRUE
## density	49.85194	66.92942	84.00690	7.866840e-15	TRUE
## pH	31.24334	43.06890	54.89446	4.727196e-13	TRUE
## sulphates	47.18117	60.72162	74.26206	7.519109e-19	TRUE
## alcohol	126.91884	165.91158	204.90433	3.729681e-17	TRUE

```
# Variable Importance Plot
plot(vi)
```



```
# Predict
# https://www.rdocumentation.org/packages/randomForestSRC/versions/3.1.0/topics/predict.rfsrc
randomForestSRC::predict.rfsrc(rf, rf_wine_test)
```

```
## Sample size of test (predict) data: 318
## Number of grow trees: 500
## Average no. of grow terminal nodes: 151.684
## Total no. of grow variables: 11
## Resampling used to grow trees: swor
## Resample size used to grow trees: 810
## Analysis: RF-R
## Family: regr
## R squared: 0.45687818
## Requested performance error: 0.54312182
```

Partial Least Squares

```
tctrl <- trainControl(method = "repeatedcv", repeats = 5, number = 10)

set.seed(4)
pls_wine <- train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data = train_transformed,
  method = "pls",
  preProc = c("center", "scale", "BoxCox"),
  tunelength = 20,
  trControl = tctrl)

pls_wine
```

```
## Partial Least Squares
##
```

```
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 1153, 1153, 1153, 1154, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      0.8009483  0.3609283  0.6248827
##   2      0.8007016  0.3612925  0.6244046
##   3      0.8008902  0.3610147  0.6242525
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 2.
```

Mars Tuning

```
mars_wine <- earth(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data=train_transformed)
```

```
mars_wine
```

```
## Selected 10 of 16 terms, and 5 of 5 predictors
## Termination condition: Reached nk 21
## Importance: alcohol, sulphates, volatile.acidity, total.sulfur.dioxide, ...
## Number of terms at each degree of interaction: 1 9 (additive model)
## GCV 0.6192871   RSS 769.9483   GRSq 0.3811964   RSq 0.3984779
```

```
summary(mars_wine)
```

```
## Call: earth(formula=quality~volatile.acidity+chlorides+total.sulfur.di...),
##           data=train_transformed)
```

```
##
##                               coefficients
## (Intercept)                   31.2282353
## h(1.77391-volatile.acidity)    0.2069420
## h(volatile.acidity-1.77391)   -0.4839817
## h(chlorides- -0.385981)       -0.0786087
## h(total.sulfur.dioxide- -1.1342) -8.4537879
## h(total.sulfur.dioxide-1.36842) -0.3247463
## h(2.48405-total.sulfur.dioxide) -8.3850064
## h(total.sulfur.dioxide-2.48405)  9.1601130
## h(0.960892-sulphates)         -0.3652400
## h(1.94253-alcohol)            -0.3674235
##
```

```
## Selected 10 of 16 terms, and 5 of 5 predictors
## Termination condition: Reached nk 21
## Importance: alcohol, sulphates, volatile.acidity, total.sulfur.dioxide, ...
## Number of terms at each degree of interaction: 1 9 (additive model)
## GCV 0.6192871   RSS 769.9483   GRSq 0.3811964   RSq 0.3984779
```



```

preProc_Arguments = c("center", "scale")
marsGrid_wine = expand.grid(.degree=1:2, .nprune=2:38)

set.seed(4)

marsModel_wine = train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data =train_transformed,
  method="earth",
  preProc=preProc_Arguments,
  tuneGrid=marsGrid_wine)

marsModel_wine

```

```

## Multivariate Adaptive Regression Spline
##
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 1281, 1281, 1281, 1281, 1281, 1281, ...
## Resampling results across tuning parameters:
##
##  degree  nprune  RMSE      Rsquared  MAE
##  1         2    0.8931576  0.2276817  0.7087136
##  1         3    0.8411903  0.3132186  0.6523206
##  1         4    0.8085427  0.3659858  0.6312428
##  1         5    0.8066254  0.3690233  0.6291302
##  1         6    0.8088540  0.3658383  0.6301515
##  1         7    0.8092640  0.3654791  0.6298044
##  1         8    0.8113355  0.3621660  0.6305667
##  1         9    0.8121219  0.3612320  0.6307932
##  1        10    0.8129649  0.3597762  0.6312176
##  1        11    0.8133643  0.3593854  0.6313570
##  1        12    0.8140167  0.3585976  0.6317295
##  1        13    0.8139652  0.3588677  0.6317555
##  1        14    0.8156289  0.3565980  0.6330332
##  1        15    0.8153264  0.3570322  0.6328053
##  1        16    0.8153866  0.3569681  0.6327920
##  1        17    0.8153866  0.3569681  0.6327920
##  1        18    0.8153866  0.3569681  0.6327920
##  1        19    0.8153866  0.3569681  0.6327920
##  1        20    0.8153866  0.3569681  0.6327920
##  1        21    0.8153866  0.3569681  0.6327920
##  1        22    0.8153866  0.3569681  0.6327920
##  1        23    0.8153866  0.3569681  0.6327920
##  1        24    0.8153866  0.3569681  0.6327920
##  1        25    0.8153866  0.3569681  0.6327920
##  1        26    0.8153866  0.3569681  0.6327920
##  1        27    0.8153866  0.3569681  0.6327920
##  1        28    0.8153866  0.3569681  0.6327920
##  1        29    0.8153866  0.3569681  0.6327920
##  1        30    0.8153866  0.3569681  0.6327920
##  1        31    0.8153866  0.3569681  0.6327920

```

```
## 1 32 0.8153866 0.3569681 0.6327920
## 1 33 0.8153866 0.3569681 0.6327920
## 1 34 0.8153866 0.3569681 0.6327920
## 1 35 0.8153866 0.3569681 0.6327920
## 1 36 0.8153866 0.3569681 0.6327920
## 1 37 0.8153866 0.3569681 0.6327920
## 1 38 0.8153866 0.3569681 0.6327920
## 2 2 0.8915976 0.2298095 0.7072556
## 2 3 0.8348954 0.3239479 0.6462899
## 2 4 0.8065778 0.3694178 0.6277609
## 2 5 0.8052118 0.3718967 0.6252944
## 2 6 0.8030613 0.3747521 0.6231283
## 2 7 0.8083906 0.3673077 0.6256714
## 2 8 0.8143151 0.3592524 0.6277289
## 2 9 0.8191785 0.3523725 0.6307293
## 2 10 0.8216029 0.3493497 0.6322984
## 2 11 0.8230969 0.3476627 0.6341852
## 2 12 0.8229519 0.3481372 0.6339090
## 2 13 0.8249943 0.3455015 0.6353646
## 2 14 0.8263816 0.3437516 0.6360541
## 2 15 0.8271766 0.3428694 0.6365014
## 2 16 0.8279797 0.3419700 0.6365412
## 2 17 0.8281539 0.3417827 0.6365777
## 2 18 0.8281539 0.3417827 0.6365777
## 2 19 0.8281539 0.3417827 0.6365777
## 2 20 0.8281539 0.3417827 0.6365777
## 2 21 0.8281539 0.3417827 0.6365777
## 2 22 0.8281539 0.3417827 0.6365777
## 2 23 0.8281539 0.3417827 0.6365777
## 2 24 0.8281539 0.3417827 0.6365777
## 2 25 0.8281539 0.3417827 0.6365777
## 2 26 0.8281539 0.3417827 0.6365777
## 2 27 0.8281539 0.3417827 0.6365777
## 2 28 0.8281539 0.3417827 0.6365777
## 2 29 0.8281539 0.3417827 0.6365777
## 2 30 0.8281539 0.3417827 0.6365777
## 2 31 0.8281539 0.3417827 0.6365777
## 2 32 0.8281539 0.3417827 0.6365777
## 2 33 0.8281539 0.3417827 0.6365777
## 2 34 0.8281539 0.3417827 0.6365777
## 2 35 0.8281539 0.3417827 0.6365777
## 2 36 0.8281539 0.3417827 0.6365777
## 2 37 0.8281539 0.3417827 0.6365777
## 2 38 0.8281539 0.3417827 0.6365777
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were nprune = 6 and degree = 2.
```

KNN Neighbors

```
set.seed(4)
```

```
knn_wine <- train(quality~ volatile.acidity + chlorides + total.sulfur.dioxide +
  sulphates + alcohol, data =train_transformed,
  method = "knn",
  preProc = c("center", "scale"),
  tuneGrid = data.frame(.k = 1:50),
  trControl = trainControl(method = "cv"))
```

```
knn_wine
```

```
## k-Nearest Neighbors
##
## 1281 samples
##    5 predictor
##
## Pre-processing: centered (5), scaled (5)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1153, 1153, 1153, 1154, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##  k    RMSE      Rsquared    MAE
##  1  0.9644161  0.2847723  0.5544215
##  2  0.8907482  0.2982523  0.6109975
##  3  0.8522988  0.3217443  0.6080274
##  4  0.8403923  0.3282242  0.6084252
##  5  0.8295624  0.3359019  0.6115016
##  6  0.8244215  0.3388616  0.6159714
##  7  0.8142260  0.3511890  0.6119213
##  8  0.8035834  0.3645164  0.6072218
##  9  0.8007209  0.3666301  0.6069772
## 10  0.7961920  0.3717519  0.6060896
## 11  0.7966255  0.3701696  0.6055852
## 12  0.7966734  0.3696525  0.6085908
## 13  0.7934338  0.3744318  0.6085134
## 14  0.7934161  0.3742458  0.6097105
## 15  0.7953655  0.3712114  0.6117051
## 16  0.7942764  0.3721523  0.6118128
## 17  0.7975641  0.3672746  0.6147680
## 18  0.7976932  0.3671614  0.6162870
## 19  0.7961296  0.3692235  0.6153625
## 20  0.7959554  0.3696504  0.6162101
## 21  0.7939922  0.3721527  0.6161404
## 22  0.7944903  0.3715377  0.6166030
## 23  0.7921240  0.3747960  0.6143588
## 24  0.7927152  0.3737572  0.6153292
## 25  0.7949951  0.3701811  0.6177137
## 26  0.7926743  0.3739579  0.6173598
## 27  0.7916786  0.3758599  0.6164459
## 28  0.7909634  0.3770866  0.6159224
## 29  0.7909416  0.3771263  0.6163367
## 30  0.7907118  0.3777643  0.6165737
## 31  0.7901812  0.3786153  0.6157317
## 32  0.7913889  0.3767080  0.6171446
## 33  0.7910876  0.3772137  0.6175664
## 34  0.7899098  0.3790728  0.6162648
```

```
## 35 0.7894065 0.3801164 0.6156748
## 36 0.7890555 0.3809280 0.6152512
## 37 0.7882048 0.3824843 0.6143572
## 38 0.7881174 0.3829143 0.6139271
## 39 0.7889098 0.3818575 0.6153589
## 40 0.7884420 0.3826805 0.6150324
## 41 0.7896330 0.3807016 0.6160041
## 42 0.7901894 0.3797496 0.6164813
## 43 0.7906953 0.3789899 0.6165895
## 44 0.7904117 0.3795530 0.6155542
## 45 0.7894754 0.3813497 0.6151459
## 46 0.7889741 0.3826017 0.6160122
## 47 0.7887943 0.3830243 0.6162245
## 48 0.7887811 0.3830731 0.6165738
## 49 0.7882914 0.3841124 0.6168664
## 50 0.7877509 0.3851382 0.6165076
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 50.
```

SVM

```
set.seed(4)

subset(train_transformed, select = -c(quality_target, quality)) -> predictors
train_transformed$quality -> quality

svmTune <- train(predictors, quality,
                 method = "svmRadial",
                 preProc = c("center", "scale"),
                 tuneLength= 5,
                 trControl = trainControl(method = "cv"))
svmTune

## Support Vector Machines with Radial Basis Function Kernel
##
## 1281 samples
## 11 predictor
##
## Pre-processing: centered (11), scaled (11)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1153, 1153, 1153, 1154, 1153, 1152, ...
## Resampling results across tuning parameters:
##
##  C      RMSE      Rsquared  MAE
##  0.25  0.7838408  0.3910168  0.5760806
##  0.50  0.7754609  0.4027211  0.5643058
##  1.00  0.7694682  0.4115229  0.5554067
##  2.00  0.7718078  0.4109309  0.5536613
##  4.00  0.7840375  0.3987760  0.5599795
##
## Tuning parameter 'sigma' was held constant at a value of 0.09900808
```

```
## RMSE was used to select the optimal model using the smallest value.  
## The final values used for the model were sigma = 0.09900808 and C = 1.
```